

Mariners Challenge

March 16, 2020

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline

import xgboost as xgb
from xgboost import XGBClassifier

from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, precision_recall_curve, auc, \
    matthews_corrcoef, accuracy_score

from joblib import load
```

```
[2]: train_df = pd.read_csv('../Data/2020-train.csv')
```

```
[3]: train_df.head()
```

```
[3]:  pitcher_id pitcher_side batter_id batter_side stadium_id umpire_id \
0    d7e3acce          Right  32678d8d          Right   a4833794  f88d09f4
1    44ec1bf5          Right   81d51733          Left   f60d6ea5  b67d862c
2    44d87ee6          Left   8eefccb7          Right   a9b8b538  13993d26
3    ff6adae0          Right   8f8ab5af          Right   e569ec39  0d8ba4bb
4    c70c96e5          Right  10874746          Right   a5ce1bf6  94a4c552

   catcher_id  inning  top_bottom  outs  ...  zone_speed  vert_approach_angle \
0    83cdf9ff        3           1  0.0  ...    86.024200             -4.37258
1    a126f66f        6           2  0.0  ...    89.458199             -4.90467
2    9db4e46f        5           2  2.0  ...    75.593597             -6.00728
3    bbbfd290        5           1  2.0  ...    76.396400             -9.50640
4    75087ec8        8           1  2.0  ...    83.215302             -4.53233

   horz_approach_angle  zone_time      x55  y55      z55  pitch_type \
0           1.429580    0.404622 -0.059343   55    6.03322          FA
1           -2.148410    0.385719 -2.148680   55    6.23380          FA
2           -0.122044    0.463953  1.300450   55    6.14750          CH
```

3	-2.581980	0.458471	-1.659590	55	6.60043	CU
4	-0.268188	0.415965	-1.526170	55	4.77332	FA

	pitch_call	pitch_id
0	InPlay	42fce2f6
1	InPlay	3e9cda86
2	BallCalled	f129a6cd
3	InPlay	03e9bc05
4	StrikeCalled	48feb675

[5 rows x 36 columns]

```
[4]: train_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 582205 entries, 0 to 582204
Data columns (total 36 columns):
#   Column                Non-Null Count  Dtype
---  -
0   pitcher_id            582205 non-null object
1   pitcher_side          582205 non-null object
2   batter_id             582205 non-null object
3   batter_side           582205 non-null object
4   stadium_id            582205 non-null object
5   umpire_id             582205 non-null object
6   catcher_id           582205 non-null object
7   inning               582205 non-null int64
8   top_bottom           582205 non-null int64
9   outs                 582053 non-null float64
10  balls                582205 non-null int64
11  strikes              582205 non-null int64
12  release_speed        582093 non-null float64
13  vert_release_angle   582093 non-null float64
14  horz_release_angle   582093 non-null float64
15  spin_rate            573194 non-null float64
16  spin_axis            582093 non-null float64
17  tilt                 580953 non-null object
18  rel_height           582093 non-null float64
19  rel_side             582093 non-null float64
20  extension            582093 non-null float64
21  vert_break           582093 non-null float64
22  induced_vert_break   582093 non-null float64
23  horz_break           582093 non-null float64
24  plate_height         582139 non-null float64
25  plate_side           582139 non-null float64
26  zone_speed           582093 non-null float64
27  vert_approach_angle  582093 non-null float64
```

```

28 horz_approach_angle 582093 non-null float64
29 zone_time           582093 non-null float64
30 x55                 582093 non-null float64
31 y55                 582205 non-null int64
32 z55                 582093 non-null float64
33 pitch_type          581720 non-null object
34 pitch_call          582205 non-null object
35 pitch_id            582205 non-null object
dtypes: float64(20), int64(5), object(11)
memory usage: 159.9+ MB

```

1 Data Cleaning

First let's get rid of the null values floating around in the data set. There are a total of 582,205 entries in the initial data set, some columns are full but some columns have null values in them. Starting with release speed, it does not seem like there are many null values, let's take a look.

```
[5]: train_df[train_df['release_speed'].isnull()]
```

```

[5]:      pitcher_id pitcher_side batter_id batter_side stadium_id umpire_id \
3405      8759809c      Right  00bde845      Right  80756f45  9c6cbb5e
12334     f6d227a5      Right  69426d29      Left   a3f610ed  9c6cbb5e
13895     b74a40d9      Right  00bde845      Right  a3f610ed  9c6cbb5e
16260     47032f76      Left   20bf9444      Right  c9712626  9c6cbb5e
17008     96a1cebe      Right  76c0475e      Right  a5ce1bf6  9c6cbb5e
...
564267    b3336756      Left  bbbfd290      Right  9b5daeaf  9c6cbb5e
564950    c3ededfb      Right  8f8ab5af      Right  b20853fa  9c6cbb5e
567947    ad6bf2a7      Right  066e327d      Right  fe6b0f40  9c6cbb5e
571936    0a606b2d      Right  336a9f05      Left   1a39a252  9c6cbb5e
573697    f556cf83      Left  30300714      Right  b20853fa  9c6cbb5e

      catcher_id  inning  top_bottom  outs  ...  zone_speed  \
3405      9c6cbb5e      5           2  0.0  ...      NaN
12334      9c6cbb5e      5           2  2.0  ...      NaN
13895      9c6cbb5e      3           1  1.0  ...      NaN
16260      9c6cbb5e     16           2  1.0  ...      NaN
17008      9c6cbb5e      9           2  0.0  ...      NaN
...
564267      9c6cbb5e      9           1  2.0  ...      NaN
564950      9c6cbb5e      1           1  2.0  ...      NaN
567947      9c6cbb5e      5           2  0.0  ...      NaN
571936      9c6cbb5e      9           1  1.0  ...      NaN
573697      9c6cbb5e      7           1  0.0  ...      NaN

      vert_approach_angle  horz_approach_angle  zone_time  x55  y55  z55  \
3405                  NaN                  NaN      NaN  NaN  55  NaN

```

12334	NaN	NaN	NaN	NaN	NaN	55	NaN
13895	NaN	NaN	NaN	NaN	NaN	55	NaN
16260	NaN	NaN	NaN	NaN	NaN	55	NaN
17008	NaN	NaN	NaN	NaN	NaN	55	NaN
...
564267	NaN	NaN	NaN	NaN	NaN	55	NaN
564950	NaN	NaN	NaN	NaN	NaN	55	NaN
567947	NaN	NaN	NaN	NaN	NaN	55	NaN
571936	NaN	NaN	NaN	NaN	NaN	55	NaN
573697	NaN	NaN	NaN	NaN	NaN	55	NaN

	pitch_type	pitch_call	pitch_id
3405	NaN	StrikeCalled	1d66612a
12334	NaN	InPlay	beb842a8
13895	NaN	InPlay	e271ef9d
16260	FA	InPlay	49f26761
17008	SL	BallCalled	ca8c6341
...
564267	CH	InPlay	8d21a585
564950	NaN	BallCalled	142a06b3
567947	NaN	StrikeCalled	5326a5f6
571936	NaN	BallCalled	3c4e2fe4
573697	NaN	StrikeCalled	763a95f4

[112 rows x 36 columns]

Only 112, I feel comfortable dropping these and it not affecting the integrity of the data.

```
[6]: train_df = train_df.drop(train_df[train_df['release_speed'].isnull()].index)
train_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 582093 entries, 0 to 582204
Data columns (total 36 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   pitcher_id            582093 non-null  object
1   pitcher_side          582093 non-null  object
2   batter_id             582093 non-null  object
3   batter_side           582093 non-null  object
4   stadium_id            582093 non-null  object
5   umpire_id             582093 non-null  object
6   catcher_id           582093 non-null  object
7   inning               582093 non-null  int64
8   top_bottom           582093 non-null  int64
9   outs                 581941 non-null  float64
10  balls                582093 non-null  int64
11  strikes              582093 non-null  int64
```

```

12 release_speed      582093 non-null float64
13 vert_release_angle 582093 non-null float64
14 horz_release_angle 582093 non-null float64
15 spin_rate          573194 non-null float64
16 spin_axis          582093 non-null float64
17 tilt               580953 non-null object
18 rel_height         582093 non-null float64
19 rel_side           582093 non-null float64
20 extension          582093 non-null float64
21 vert_break         582093 non-null float64
22 induced_vert_break 582093 non-null float64
23 horz_break         582093 non-null float64
24 plate_height       582093 non-null float64
25 plate_side         582093 non-null float64
26 zone_speed         582093 non-null float64
27 vert_approach_angle 582093 non-null float64
28 horz_approach_angle 582093 non-null float64
29 zone_time          582093 non-null float64
30 x55                582093 non-null float64
31 y55                582093 non-null int64
32 z55                582093 non-null float64
33 pitch_type         581674 non-null object
34 pitch_call         582093 non-null object
35 pitch_id           582093 non-null object
dtypes: float64(20), int64(5), object(11)
memory usage: 164.3+ MB

```

Now let's take a look at outs. Again, not many null values in it, let's see how many there are.

```
[7]: train_df[train_df['outs'].isnull()]
```

```

[7]:      pitcher_id pitcher_side batter_id batter_side stadium_id umpire_id \
6689      a6118212      Right  5dce2d1c      Right   99faafae  c16da957
14410      5b740fab      Left  e2e2a336      Right   d0e0eb76  c9752165
19161      161160dd      Right f338e9d3      Left    1a39a252  c229ef9e
23843      f6d227a5      Right e9553a98      Right   d0e0eb76  26a1bb6b
27156      264562c6      Right 29d12af7      Left    a5ce1bf6  9806dfbc
...      ...      ...      ...      ...      ...
559476      b3d5c0a9      Right cf690f2f      Left    0c59f5af  8f1ef267
569804      91700130      Right de9d396f      Left    99faafae  667d5752
569836      91700130      Right de9d396f      Left    99faafae  667d5752
570351      d5ef78cb      Right 781ec6be      Right   aa998b21  a9ad7586
574039      b3d5c0a9      Right cf690f2f      Left    0c59f5af  8f1ef267

      catcher_id  inning  top_bottom  outs  ...  zone_speed  \
6689      fd37f21c      7           1   NaN  ...   75.911797
14410      5a42193e      5           2   NaN  ...   72.887398
19161      41ac8158      4           2   NaN  ...   84.766296

```

23843	Offec018	6	1	NaN	...	74.321999
27156	a3a2988b	11	2	NaN	...	86.122299
...
559476	fa18ff59	5	2	NaN	...	80.375298
569804	41ac8158	7	1	NaN	...	84.243103
569836	41ac8158	7	1	NaN	...	82.846298
570351	a3a2988b	6	1	NaN	...	85.266098
574039	fa18ff59	5	2	NaN	...	67.233902

	vert_approach_angle	horz_approach_angle	zone_time	x55	y55	\
6689	-10.27540	-1.215380	0.457509	-1.44079	55	
14410	-7.79636	0.764407	0.476133	3.30888	55	
19161	-5.21754	0.053372	0.417402	-1.30256	55	
23843	-9.18817	-4.715870	0.470084	-2.66638	55	
27156	-7.29354	-4.722650	0.411571	-2.93787	55	
...
559476	-7.07432	-2.223690	0.430519	-1.15777	55	
569804	-5.38618	0.594491	0.419565	-1.64532	55	
569836	-6.36660	-1.029200	0.425654	-1.62191	55	
570351	-4.60246	-1.696110	0.401914	-1.92533	55	
574039	-11.05840	-2.251130	0.515505	-1.19462	55	

	z55	pitch_type	pitch_call	pitch_id
6689	6.19782	SL	StrikeSwinging	8f7e287a
14410	5.70064	CH	BallCalled	fdaab946
19161	5.76697	FA	BallCalled	479ee407
23843	6.70092	CH	StrikeSwinging	b01234ef
27156	6.37265	SL	StrikeSwinging	f7771432
...
559476	6.08150	SL	FoulBall	d29aa13d
569804	6.31180	FA	BallCalled	d711100f
569836	6.24978	FA	StrikeCalled	5f05e674
570351	5.80506	FA	StrikeSwinging	24ac5f78
574039	5.88421	CU	FoulBall	3a28f2f5

[152 rows x 36 columns]

Only 152, that's good. I'll drop these as well since there's no real way to know how many outs there were on a given pitch.

```
[8]: train_df = train_df.drop(train_df[train_df['outs'].isnull()].index)
train_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 581941 entries, 0 to 582204
Data columns (total 36 columns):
#   Column              Non-Null Count  Dtype
---
```

```

0    pitcher_id      581941 non-null object
1    pitcher_side    581941 non-null object
2    batter_id       581941 non-null object
3    batter_side     581941 non-null object
4    stadium_id      581941 non-null object
5    umpire_id       581941 non-null object
6    catcher_id      581941 non-null object
7    inning          581941 non-null int64
8    top_bottom      581941 non-null int64
9    outs            581941 non-null float64
10   balls           581941 non-null int64
11   strikes         581941 non-null int64
12   release_speed   581941 non-null float64
13   vert_release_angle 581941 non-null float64
14   horz_release_angle 581941 non-null float64
15   spin_rate       573044 non-null float64
16   spin_axis       581941 non-null float64
17   tilt            580802 non-null object
18   rel_height      581941 non-null float64
19   rel_side        581941 non-null float64
20   extension       581941 non-null float64
21   vert_break      581941 non-null float64
22   induced_vert_break 581941 non-null float64
23   horz_break      581941 non-null float64
24   plate_height    581941 non-null float64
25   plate_side      581941 non-null float64
26   zone_speed      581941 non-null float64
27   vert_approach_angle 581941 non-null float64
28   horz_approach_angle 581941 non-null float64
29   zone_time       581941 non-null float64
30   x55             581941 non-null float64
31   y55             581941 non-null int64
32   z55             581941 non-null float64
33   pitch_type      581522 non-null object
34   pitch_call      581941 non-null object
35   pitch_id        581941 non-null object
dtypes: float64(20), int64(5), object(11)
memory usage: 164.3+ MB

```

Now let's take a look at spin rate. It certainly looks like there are a lot of data points missing in this column, so dropping all the values may not be the best idea. Let's see how many there are.

```
[9]: train_df[train_df['spin_rate'].isnull()]
```

```

[9]:    pitcher_id pitcher_side batter_id batter_side stadium_id umpire_id \
43      4c807a49         Left  210e8d5b         Right  402559d3  4ff102e5
148     cb113772         Right  96339e13         Left  aa998b21  c683b9a6
161     af6d3149         Right  6c43d395         Right  03722f5d  d057fd71

```

237	eccb6087	Right	08b0b39d	Left	a5ce1bf6	9806dfbc
304	e332e67d	Right	34a8f234	Left	402559d3	fbbea103
...
581856	00f5fb90	Right	b4efd4bf	Right	f682daed	cac8185e
581859	fa9b0925	Right	c1ec06e6	Right	5025d8df	7675ce83
581927	09da5d7a	Left	073c2b16	Right	b20853fa	4db7bcb
581992	be5181f0	Right	566220c7	Right	03722f5d	a86853a2
582189	a2f05755	Right	e7a70ed1	Left	854c6c72	16750c18

	catcher_id	inning	top_bottom	outs	...	zone_speed	\
43	fb0970f	6	2	2.0	...	77.869301	
148	a3a2988b	7	1	0.0	...	73.582497	
161	e9aa50df	9	1	1.0	...	76.742798	
237	b1499101	9	1	1.0	...	79.280602	
304	e4fac104	5	2	2.0	...	80.649803	
...	
581856	e4fac104	4	2	1.0	...	77.819801	
581859	054f7d9f	6	1	0.0	...	77.702797	
581927	9db4e46f	9	1	0.0	...	76.172997	
581992	5b8927f6	2	2	2.0	...	77.229103	
582189	daa1322d	4	2	2.0	...	77.375603	


	vert_approach_angle	horz_approach_angle	zone_time	x55	y55	\
43	-6.96547	2.47110	0.444589	2.32704	55	
148	-9.61581	-2.45899	0.474079	-1.42864	55	
161	-6.86927	-3.62681	0.461172	-3.17556	55	
237	-6.75395	-3.49224	0.438489	-1.96139	55	
304	-8.90885	-2.33627	0.433380	-1.18638	55	
...	
581856	-8.53157	-2.76430	0.446821	-2.01818	55	
581859	-10.84580	-3.94847	0.452332	-1.45562	55	
581927	-8.79066	3.64529	0.459917	2.00739	55	
581992	-7.42066	-3.82928	0.458138	-1.49327	55	
582189	-8.75378	-2.46304	0.450004	-1.84355	55	

	z55	pitch_type	pitch_call	pitch_id
43	5.75515	SL	InPlay	1769b4d5
148	5.84820	SL	BallCalled	ef2a6b0e
161	4.88357	SL	StrikeCalled	f31fc865
237	5.75645	SL	StrikeCalled	3588e119
304	6.42062	SL	InPlay	6fa4b697
...
581856	5.79492	SL	InPlay	fc151903
581859	6.02164	CU	BallCalled	941e91ca
581927	6.51264	SL	StrikeSwinging	3f7dd184
581992	5.40118	SL	StrikeSwinging	40cc471a
582189	5.49470	SL	InPlay	059d09ff

[8897 rows x 36 columns]

Almost 9000 rows. Dropping that many data points could have an adverse effect on the data set and any potential modeling we do. Let's take a look at a countplot and see the distribution of the spin rate column.

```
[10]: plt.figure(figsize=(12,9))
sns.distplot(train_df[train_df['spin_rate'].notnull()]['spin_rate'], hist=True,
             kde=False)
plt.xlabel('Spin rate')
plt.ylabel('Count')
plt.title('Count of spin rates')
plt.show()
```



It looks like a normal shaped curve, which is expected from this type of data. Filling the missing values with the average would be the best way to go. That keeps the normal distribution of this variable intact and shouldn't have any adverse effect on any future modeling.

```
[11]: train_df['spin_rate'].describe()
```

```
[11]: count    573044.000000
      mean      2220.693335
      std       311.989506
      min       362.382996
      25%       2072.879883
      50%       2238.449951
      75%       2400.790039
      max       3752.239990
      Name: spin_rate, dtype: float64
```

```
[12]: train_df['spin_rate'].fillna(train_df['spin_rate'].mean(), inplace=True)
```

```
[13]: train_df['spin_rate'].describe()
```

```
[13]: count    581941.000000
      mean      2220.693335
      std       309.595392
      min       362.382996
      25%       2076.080078
      50%       2233.979980
      75%       2397.649902
      max       3752.239990
      Name: spin_rate, dtype: float64
```

Tilt is the next column to have missing values, let's take a look at how many.

```
[14]: train_df[train_df['tilt'].isnull()]
```

```
[14]:
```

	pitcher_id	pitcher_side	batter_id	batter_side	stadium_id	umpire_id	\
378	cd483725	Right	192899a6	Right	c9712626	1869cf54	
770	22b76a09	Left	d11080ae	Left	5025d8df	598ea1b2	
1339	cd483725	Right	deb2ab32	Right	0b15e1ca	eb059a22	
2117	28e273c4	Left	4ac005f3	Right	d0d69f32	46051258	
3399	193d153f	Left	ed874f19	Right	c9712626	eb059a22	
...	
580815	98eaf8b2	Right	f57085ec	Right	99faafae	373947e5	
581637	44ec1bf5	R	b3dac04c	L	cfe02944	aea4dd5a	
581716	60a6f8df	Right	cddcbd8f	Left	0faa3b2d	4581c636	
581733	f45c0602	Right	ad84b429	Left	fe6b0f40	bb04ea23	
582161	1fb18290	Left	e14059d7	Right	6a69d99b	ff7406e8	

	catcher_id	inning	top_bottom	outs	...	zone_speed	\
378	00ae6fb5	3	1	1.0	...	73.954399	
770	a421b54b	3	2	0.0	...	67.967903	
1339	00ae6fb5	6	2	1.0	...	75.172897	
2117	4fedda83	3	1	1.0	...	75.375900	
3399	4f9cd7f9	4	1	2.0	...	75.718102	

```

...
580815  ccd72da8      9      2  2.0 ... 79.889999
581637  a126f66f      6      2  1.0 ... 88.644798
581716  a421b54b      7      1  0.0 ... 79.617599
581733  dc18f830      1      1  2.0 ... 82.575500
582161  bbbfd290      8      2  2.0 ... 71.934700

      vert_approach_angle  horz_approach_angle  zone_time      x55  y55  \
378          -9.60881      -1.961030    0.477100 -1.181070  55
770         -11.73220       1.604870    0.512340  1.140470  55
1339        -9.79158      -1.867760    0.471512 -0.749778  55
2117        -8.49879       2.576980    0.459569  2.623400  55
3399        -9.39808       2.573850    0.462214  2.150780  55
...
580815          -8.93767      -1.194220    0.435881 -1.759100  55
581637          -5.31458      -2.812660    0.390148 -2.398420  55
581716          -6.18068      -2.975590    0.435664 -2.043060  55
581733          -5.96087      -0.656144    0.416289 -2.912520  55
582161         -13.77550       1.224220    0.486536  1.447500  55

      z55  pitch_type      pitch_call  pitch_id
378    6.13692      CU      InPlay  d30f5214
770    6.32628      CU    BallCalled  a7209bc8
1339    5.86303      CU  StrikeSwinging  6b05ddc4
2117    6.40245      CU      InPlay  3face29a
3399    6.07629      CU  StrikeSwinging  e1c3703d
...
580815    6.48687      CU  StrikeSwinging  672aa57b
581637    6.25742      FA    BallCalled  9e8fb97a
581716    5.01476      SL    StrikeCalled  1d36d6f5
581733    5.64645      CH      FoulBall  2e8f759b
582161    6.16835      CU  StrikeSwinging  7db0fd25

```

[1139 rows x 36 columns]

1139 rows with null values. That's not much in the grand scheme of the data set, only about 0.2% of the entire data set. Dropping these would not hurt in the long run.

```
[15]: train_df = train_df.drop(train_df[train_df['tilt'].isnull()].index)
```

```
[16]: train_df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 580802 entries, 0 to 582204
Data columns (total 36 columns):
#   Column              Non-Null Count  Dtype
---  -
0   pitcher_id          580802 non-null  object

```

```

1  pitcher_side      580802 non-null object
2  batter_id        580802 non-null object
3  batter_side      580802 non-null object
4  stadium_id       580802 non-null object
5  umpire_id        580802 non-null object
6  catcher_id       580802 non-null object
7  inning           580802 non-null int64
8  top_bottom       580802 non-null int64
9  outs             580802 non-null float64
10 balls            580802 non-null int64
11 strikes          580802 non-null int64
12 release_speed    580802 non-null float64
13 vert_release_angle 580802 non-null float64
14 horz_release_angle 580802 non-null float64
15 spin_rate        580802 non-null float64
16 spin_axis        580802 non-null float64
17 tilt            580802 non-null object
18 rel_height       580802 non-null float64
19 rel_side         580802 non-null float64
20 extension        580802 non-null float64
21 vert_break       580802 non-null float64
22 induced_vert_break 580802 non-null float64
23 horz_break       580802 non-null float64
24 plate_height     580802 non-null float64
25 plate_side       580802 non-null float64
26 zone_speed       580802 non-null float64
27 vert_approach_angle 580802 non-null float64
28 horz_approach_angle 580802 non-null float64
29 zone_time        580802 non-null float64
30 x55              580802 non-null float64
31 y55              580802 non-null int64
32 z55              580802 non-null float64
33 pitch_type       580383 non-null object
34 pitch_call       580802 non-null object
35 pitch_id         580802 non-null object
dtypes: float64(20), int64(5), object(11)
memory usage: 164.0+ MB

```

Finally, `pitch_type` is the last column to have null values in it. Let's take a look.

```
[17]: train_df[train_df['pitch_type'].isnull()]
```

```

[17]:   pitcher_id pitcher_side batter_id batter_side stadium_id umpire_id \
634      bff0f759      Left  192899a6      Right   d0d69f32  9c6cbb5e
3126     900e6090     Right  699983d6      Left    78aaa563  9c6cbb5e
3263     7bdd4794     Right  44924919     Right    83508f28  9c6cbb5e
3485     7bdd4794     Right  0b8c61b3      Left    83508f28  9c6cbb5e
4655     57613174     Right  00bde845     Right    78aaa563  9c6cbb5e

```

...
576312	b48cf592	Right	a3b17b9b	Left	cfe02944	9c6cbb5e
578590	b4eadd6d	Right	0ae0de45	Right	0a0cfe0d	9c6cbb5e
579026	57613174	Right	6b115fe9	Left	fe6b0f40	9c6cbb5e
579999	d629b647	Right	699983d6	Left	78aaa563	9c6cbb5e
580965	57613174	Right	fd347bb1	Left	78aaa563	9c6cbb5e

	catcher_id	inning	top_bottom	outs	...	zone_speed	\
634	9c6cbb5e	5	2	0.0	...	67.892998	
3126	9c6cbb5e	8	2	1.0	...	81.340797	
3263	9c6cbb5e	9	1	0.0	...	86.714600	
3485	9c6cbb5e	8	1	1.0	...	85.167503	
4655	9c6cbb5e	12	1	1.0	...	85.837502	

...
576312	9c6cbb5e	8	2	1.0	...	74.434799
578590	9c6cbb5e	8	1	1.0	...	87.362602
579026	9c6cbb5e	8	2	0.0	...	77.284798
579999	9c6cbb5e	14	2	2.0	...	83.999802
580965	9c6cbb5e	14	1	0.0	...	86.906998

	vert_approach_angle	horz_approach_angle	zone_time	x55	y55	\
634	-7.08524	0.721203	0.516315	2.17234	55	
3126	-5.03602	-2.062410	0.432663	-2.91834	55	
3263	-3.29493	-1.750810	0.402108	-1.73126	55	
3485	-4.46310	-1.934850	0.410087	-1.83128	55	
4655	-4.32100	-1.389630	0.402448	-1.74871	55	
...	
576312	-1.24151	-1.315220	0.465865	-1.59554	55	
578590	-4.12619	-1.554950	0.396194	-1.33279	55	
579026	-7.36536	-1.514690	0.460128	-1.91246	55	
579999	-5.26501	0.381528	0.411825	-1.65040	55	
580965	-3.75225	-0.884140	0.397061	-1.94958	55	

	z55	pitch_type	pitch_call	pitch_id
634	6.50402	NaN	StrikeCalled	34257ee5
3126	5.70209	NaN	InPlay	972b06d9
3263	5.63128	NaN	BallCalled	44187be2
3485	5.61829	NaN	StrikeCalled	bf482788
4655	5.99225	NaN	InPlay	f41ccb06
...
576312	1.38558	NaN	FoulBall	17b5318a
578590	6.22610	NaN	FoulBall	d20f5e52
579026	5.90440	NaN	FoulBall	528ecba1
579999	6.96849	NaN	FoulBall	e0d8dea9
580965	5.92467	NaN	InPlay	fba61f04

[419 rows x 36 columns]

Only 419 rows, these can be dropped.

```
[18]: train_df = train_df.drop(train_df[train_df['pitch_type'].isnull()].index)
train_df = train_df.reset_index().drop('index', axis=1)
```

```
[19]: train_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 580383 entries, 0 to 580382
Data columns (total 36 columns):
#   Column                Non-Null Count  Dtype
---  -
0   pitcher_id            580383 non-null  object
1   pitcher_side          580383 non-null  object
2   batter_id            580383 non-null  object
3   batter_side          580383 non-null  object
4   stadium_id           580383 non-null  object
5   umpire_id            580383 non-null  object
6   catcher_id           580383 non-null  object
7   inning               580383 non-null  int64
8   top_bottom           580383 non-null  int64
9   outs                 580383 non-null  float64
10  balls                580383 non-null  int64
11  strikes              580383 non-null  int64
12  release_speed        580383 non-null  float64
13  vert_release_angle   580383 non-null  float64
14  horz_release_angle   580383 non-null  float64
15  spin_rate            580383 non-null  float64
16  spin_axis            580383 non-null  float64
17  tilt                 580383 non-null  object
18  rel_height           580383 non-null  float64
19  rel_side             580383 non-null  float64
20  extension            580383 non-null  float64
21  vert_break           580383 non-null  float64
22  induced_vert_break   580383 non-null  float64
23  horz_break           580383 non-null  float64
24  plate_height         580383 non-null  float64
25  plate_side           580383 non-null  float64
26  zone_speed           580383 non-null  float64
27  vert_approach_angle  580383 non-null  float64
28  horz_approach_angle  580383 non-null  float64
29  zone_time            580383 non-null  float64
30  x55                  580383 non-null  float64
31  y55                  580383 non-null  int64
32  z55                  580383 non-null  float64
33  pitch_type           580383 non-null  object
34  pitch_call           580383 non-null  object
35  pitch_id             580383 non-null  object
```

```
dtypes: float64(20), int64(5), object(11)
memory usage: 159.4+ MB
```

```
[20]: train_df.head(10)
```

```
[20]:  pitcher_id pitcher_side batter_id batter_side stadium_id umpire_id \
0    d7e3acce          Right  32678d8d          Right   a4833794   f88d09f4
1    44ec1bf5          Right   81d51733          Left   f60d6ea5   b67d862c
2    44d87ee6          Left   8eefccb7          Right   a9b8b538   13993d26
3    ff6adae0          Right   8f8ab5af          Right   e569ec39   0d8ba4bb
4    c70c96e5          Right  10874746          Right   a5ce1bf6   94a4c552
5    98f8936a          Right  a58e31f3          Right   9b5daeaf   0dad94e8
6    28e273c4          Left   9a2db1f2          Right   d0d69f32   caf1f50b
7    4f3062b6          Left   7e2bb9a9          Right   c9712626   33bb973b
8    afae9816          Left   ffe7832e          Left    d0d69f32   f88d09f4
9    61ab8c67          Right  daa1322d          Right   f682daed   c4c41d26

    catcher_id  inning  top_bottom  outs  ...  zone_speed  vert_approach_angle \
0    83cdf9ff         3           1  0.0  ...   86.024200           -4.37258
1    a126f66f         6           2  0.0  ...   89.458199           -4.90467
2    9db4e46f         5           2  2.0  ...   75.593597           -6.00728
3    bbbfd290         5           1  2.0  ...   76.396400           -9.50640
4    75087ec8         8           1  2.0  ...   83.215302           -4.53233
5    68d1111a         7           2  0.0  ...   80.265404           -8.24794
6    4fedda83         3           1  1.0  ...   87.948799           -4.76645
7    20bf9444         6           1  1.0  ...   76.352798          -10.25710
8    4fedda83         2           1  1.0  ...   78.281097           -4.85101
9    41ac8158         4           1  1.0  ...   86.078400           -6.09955

    horz_approach_angle  zone_time    x55  y55    z55  pitch_type \
0           1.429580    0.404622 -0.059343   55  6.03322         FA
1           -2.148410    0.385719 -2.148680   55  6.23380         FA
2           -0.122044    0.463953  1.300450   55  6.14750         CH
3           -2.581980    0.458471 -1.659590   55  6.60043         CU
4           -0.268188    0.415965 -1.526170   55  4.77332         FA
5           0.780148    0.438111 -2.075230   55  5.79080         CH
6           0.696210    0.390590  2.569990   55  6.09316         FA
7           4.681720    0.464021  2.497290   55  6.22659         SL
8           1.945590    0.440157  2.559950   55  5.91159         FA
9           0.425454    0.411268 -0.876224   55  6.53540         FA

    pitch_call  pitch_id
0      InPlay  42fce2f6
1      InPlay  3e9cda86
2  BallCalled  f129a6cd
3      InPlay  03e9bc05
4  StrikeCalled  48feb675
```

```

5 StrikeSwinging 419540c7
6 StrikeSwinging cf85249d
7     BallCalled c9423da3
8     FoulBall 51ad39b4
9     FoulBall b89e4ec3

```

```
[10 rows x 36 columns]
```

No more null values in the data set now, now we can move on to creating the target variable for modeling.

2 Data Wrangling

In the test set, the target variable is called “is_strike”, and we don’t have a column like that here in the training set. However, we do have a “pitch_call” column, which we can use to create the “is_strike” column. Along with called strikes and swinging strikes, any ball batted in to play or any foul balls are also counted as strikes. Using this, we can build the “is_strike” column using a simple for loop.

```
[21]: train_df['pitch_call'].unique()
```

```
[21]: array(['InPlay', 'BallCalled', 'StrikeCalled', 'StrikeSwinging',
          'FoulBall', 'HitByPitch', 'BallIntentional'], dtype=object)
```

```
[22]: is_strike_list = ['InPlay', 'StrikeCalled', 'StrikeSwinging', 'FoulBall']
train_df['is_strike'] = train_df['pitch_call'].apply(lambda x: 1 if x in is_strike_list else 0)
```

We can easily assign that list to a new column in the data set.

```
[23]: train_df.head(10)
```

```
[23]:
```

	pitcher_id	pitcher_side	batter_id	batter_side	stadium_id	umpire_id	\
0	d7e3acce	Right	32678d8d	Right	a4833794	f88d09f4	
1	44ec1bf5	Right	81d51733	Left	f60d6ea5	b67d862c	
2	44d87ee6	Left	8eefccb7	Right	a9b8b538	13993d26	
3	ff6adae0	Right	8f8ab5af	Right	e569ec39	0d8ba4bb	
4	c70c96e5	Right	10874746	Right	a5ce1bf6	94a4c552	
5	98f8936a	Right	a58e31f3	Right	9b5daeaf	0dad94e8	
6	28e273c4	Left	9a2db1f2	Right	d0d69f32	caf1f50b	
7	4f3062b6	Left	7e2bb9a9	Right	c9712626	33bb973b	
8	afae9816	Left	ffe7832e	Left	d0d69f32	f88d09f4	
9	61ab8c67	Right	daa1322d	Right	f682daed	c4c41d26	

	catcher_id	inning	top_bottom	outs	...	vert_approach_angle	\
0	83cdf9ff	3	1	0.0	...	-4.37258	
1	a126f66f	6	2	0.0	...	-4.90467	

2	9db4e46f	5	2	2.0	...	-6.00728
3	bbbfd290	5	1	2.0	...	-9.50640
4	75087ec8	8	1	2.0	...	-4.53233
5	68d1111a	7	2	0.0	...	-8.24794
6	4fedda83	3	1	1.0	...	-4.76645
7	20bf9444	6	1	1.0	...	-10.25710
8	4fedda83	2	1	1.0	...	-4.85101
9	41ac8158	4	1	1.0	...	-6.09955

	horz_approach_angle	zone_time	x55	y55	z55	pitch_type \
0	1.429580	0.404622	-0.059343	55	6.03322	FA
1	-2.148410	0.385719	-2.148680	55	6.23380	FA
2	-0.122044	0.463953	1.300450	55	6.14750	CH
3	-2.581980	0.458471	-1.659590	55	6.60043	CU
4	-0.268188	0.415965	-1.526170	55	4.77332	FA
5	0.780148	0.438111	-2.075230	55	5.79080	CH
6	0.696210	0.390590	2.569990	55	6.09316	FA
7	4.681720	0.464021	2.497290	55	6.22659	SL
8	1.945590	0.440157	2.559950	55	5.91159	FA
9	0.425454	0.411268	-0.876224	55	6.53540	FA

	pitch_call	pitch_id	is_strike
0	InPlay	42fce2f6	1
1	InPlay	3e9cda86	1
2	BallCalled	f129a6cd	0
3	InPlay	03e9bc05	1
4	StrikeCalled	48feb675	1
5	StrikeSwinging	419540c7	1
6	StrikeSwinging	cf85249d	1
7	BallCalled	c9423da3	0
8	FoulBall	51ad39b4	1
9	FoulBall	b89e4ec3	1

[10 rows x 37 columns]

Now we have our data set with the target variable, let's take a look at the "is_strike" column and its value counts.

```
[24]: train_df['is_strike'].value_counts()
```

```
[24]: 1    369807
      0    210576
      Name: is_strike, dtype: int64
```

Interesting. We have an imbalanced classification problem here, with the majority class being almost twice as large as the minority class. That has implications for modeling in the future, namely being careful about what classification model is used for this problem. We also may need to use some resampling methods if the model is choosing the majority class by an overwhelming

margin.

Before we get into modeling however, there was something I noticed with the “tilt” column. It has two different types of string data packed into the column. We’ll need to fix that column to get it all into one data format.

```
[25]: train_df['tilt'].unique()
```

```
[25]: array(['1:00', '12:15', '11:15', '7:45', '2:15', '2:45', '10:30', '4:45',  
          '11:00', '1:30', '1:15', '12:45', '5:15', '10:45', '6:30', '8:00',  
          '4:00', '2:00', '7:30', '3:30', '12:00', '1:45', '9:00', '10:00',  
          '11:30', '32400 secs', '12:30', '9:15', '11:45', '9:45', '10:15',  
          '3:00', '42300 secs', '5:00', '7:15', '7:00', '6:45', '9:30',  
          '8:45', '3:45', '43200 secs', '6:15', '2:30', '4:15', '35100 secs',  
          '5:45', '5:30', '8:15', '8:30', '29700 secs', '44100 secs', '3:15',  
          '4:30', '14400 secs', '6300 secs', '45000 secs', '38700 secs',  
          '34200 secs', '36000 secs', '45900 secs', '36900 secs',  
          '3600 secs', '18000 secs', '7200 secs', '40500 secs', '5400 secs',  
          '8100 secs', '15300 secs', '27900 secs', '4500 secs', '23400 secs',  
          '25200 secs', '41400 secs', '30600 secs', '6:00', '9900 secs',  
          '33300 secs', '37800 secs', '13500 secs', '27000 secs',  
          '39600 secs', '12600 secs', '17100 secs', '16200 secs',  
          '11700 secs', '9000 secs', '18900 secs', '26100 secs',  
          '22500 secs', '20700 secs', '24300 secs', '31500 secs',  
          '21600 secs', '10800 secs', '19800 secs', '28800 secs'],  
        dtype=object)
```

I’m choosing to turn all of the “1:00”, “12:15” format into a seconds-from-midnight integer, that will be the easiest way to get all of the column into one data format and data type.

```
[26]: train_df['tilt'] = train_df['tilt'].map(lambda x: sum(a*int(t) for a, t in  
          ↪zip([3600, 60], x.split(':')))) \n  
          if ':' in x else int(x[:-5]))
```

```
[27]: train_df.to_csv('../Data/model_data.csv')
```

3 Predictive Modeling

Now that our data set clean and how we want it, we can get into some predictive modeling. Seeing as this is a binary classification problem, we’ll need to use a classification algorithm. I’m choosing to use gradient boosting here because I’ve used it in the past and have gotten good results with it in a timely manner. Logistic regression would be faster, but would give us a less accurate model than a gradient boosting model.

I’m going to do the hyperparameter tuning in its own dedicated notebook, then load the trained model into this notebook after it’s been fitted with all the correct hyperparameters.

```
[28]: X = train_df.drop(['pitcher_id', 'batter_id', 'stadium_id', 'umpire_id',
↳ 'catcher_id', 'pitch_call', 'is_strike', 'pitch_id'], axis=1)
y = train_df['is_strike']

X_train, X_test, y_train, y_test = train_test_split(X, y, stratify=y,
↳ random_state=34)
```

Getting dummy variables for some of the categorical variables would be good for modeling. It will allow us to see how much importance was placed on these features by the model.

```
[29]: X_train = pd.get_dummies(X_train, prefix=['pitcher', 'batter', 'is'],
↳ columns=['pitcher_side', 'batter_side', 'pitch_type'])
X_test = pd.get_dummies(X_test, prefix=['pitcher', 'batter', 'is'],
↳ columns=['pitcher_side', 'batter_side', 'pitch_type'])
```

```
[30]: X_train.head()
```

```
[30]:
```

	inning	top_bottom	outs	balls	strikes	release_speed	\
362341	8	1	1.0	3	2	86.056702	
392255	3	2	0.0	0	1	93.787697	
520345	3	1	0.0	0	0	90.838699	
120374	3	1	1.0	1	1	89.893600	
194486	7	1	1.0	0	0	94.105202	

	vert_release_angle	horz_release_angle	spin_rate	spin_axis	...	\
362341	-0.754319	-2.89459	2425.870117	205.197006	...	
392255	-1.202140	-3.24475	2496.909912	210.024994	...	
520345	-1.730600	1.75061	2230.830078	158.481995	...	
120374	-2.494800	-3.26674	2175.179932	-150.666000	...	
194486	-2.047340	-2.19327	2071.290039	227.212997	...	

	pitcher_Left	pitcher_Right	batter_Left	batter_Right	is_CH	is_CU	\
362341	0	1	0	1	0	0	
392255	0	1	1	0	0	0	
520345	1	0	0	1	0	0	
120374	0	1	0	1	0	0	
194486	0	1	1	0	0	0	

	is_FA	is_KN	is_SL	is_XX
362341	0	0	1	0
392255	1	0	0	0
520345	1	0	0	0
120374	0	0	1	0
194486	1	0	0	0

[5 rows x 36 columns]

Now that we have our modeling data set, let's get into modeling the data, first we will look at a

non-tuned XGBoost model and see how it performs on the data, then compare that to the tuned model from the “Mariners Machine Learning Model” notebook.

```
[ ]: base_xgb = XGBClassifier()
base_xgb.fit(X_train, y_train)
xgb_base_pred = base_xgb.predict(X_test)

[ ]: print(f'Base XGB Classifier Test Accuracy: {round(accuracy_score(y_test,
    ↳xgb_base_pred) * 100, 2)}')
print('Base XGB Classifier Classification Report')
print(classification_report(y_test, xgb_base_pred))
print('\n')
print(f'Base XGB Classifier MCC Score: {matthews_corrcoef(y_test,
    ↳xgb_base_pred)}')

xgb_probs = base_xgb.predict_proba(X_test)
xgb_probs = xgb_probs[:, 1]

xgb_precision, xgb_recall, _ = precision_recall_curve(y_test, xgb_probs)

no_skill = len(y_test[y_test == 1]) / len(y_test)
plt.figure(figsize=(12,9))
plt.plot([0, 1], [no_skill, no_skill], linestyle='--', label='No Skill')
plt.plot(xgb_recall, xgb_precision, marker='.', label='Gradient Boosting')
plt.title('Precision Recall Curve for a Base XGB Classifier')
plt.xlabel('Recall')
plt.ylabel('Precision')
plt.legend()
plt.show()

print(f'Base XGB Classifier AUC Score: {auc(xgb_recall, xgb_precision)}')
```

By itself, it's a good model. 86% accuracy and a solid precision/recall on both classes is a good start. As well as a .706 MCC score and very high AUC score. Looks like the base XGBoost is a good starting out point for the hyperparameter tuning we did in the other notebook. Let's load that in and take a look at the same metrics as above.

```
[ ]: tuned_xgb = load('xgboost_model.pkl')
xgb_tuned_pred = tuned_xgb.predict(X_test)

[ ]: print(f'Tuned XGB Classifier Test Accuracy: {round(accuracy_score(y_test,
    ↳xgb_tuned_pred) * 100, 2)}')
print('Tuned XGB Classifier Classification Report')
print(classification_report(y_test, xgb_tuned_pred))
print('\n')
print(f'Tuned XGB Classifier MCC Score: {matthews_corrcoef(y_test,
    ↳xgb_tuned_pred)}')
```

```

xgb_probs = tuned_xgb.predict_proba(X_test)
xgb_probs = xgb_probs[:, 1]

xgb_precision, xgb_recall, _ = precision_recall_curve(y_test, xgb_probs)

no_skill = len(y_test[y_test == 1]) / len(y_test)
plt.figure(figsize=(12,9))
plt.plot([0, 1], [no_skill, no_skill], linestyle='--', label='No Skill')
plt.plot(xgb_recall, xgb_precision, marker='.', label='Gradient Boosting')
plt.title('Precision Recall Curve for a Tuned XGB Classifier')
plt.xlabel('Recall')
plt.ylabel('Precision')
plt.legend()
plt.show()

print(f'Tuned XGB Classifier AUC Score: {auc(xgb_recall, xgb_precision)}')

```

It did better, even if it was only slightly. Precision and recall scores on both classes improved, accuracy went up, and MCC and AUC scores went up as well. The hyper-parameter tuning we did in the other notebook worked well.

3.1 Making predictions on test set

We have a good model trained, now we can make predictions on the testing set. First we need to load it in and clean it the way we cleaned the training set.

```

[ ]: test_df = pd.read_csv('../Data/2020-test.csv')

[ ]: def clean_and_wrangle(df):

    df = df.drop(df[df['release_speed'].isnull()].index)
    df = df.drop(df[df['outs'].isnull()].index)
    df['spin_rate'].fillna(df['spin_rate'].mean(), inplace=True)
    df = df.drop(df[df['tilt'].isnull()].index)
    df = df.drop(df[df['pitch_type'].isnull()].index)
    df = df.reset_index().drop('index', axis=1)

    df['tilt'] = df['tilt'].map(lambda x: sum(a*int(t) for a, t in zip([3600, 60],
↪x.split(':')))) \
                                if ':' in x else int(x[:-5]))

    return df

[ ]: test_df = clean_and_wrangle(test_df)

[ ]: test_df.head(10)

[ ]: test_df.info()

```

No null values (except our target variable “is_strike”). We’ll use our trained model from up above and make predictions on the entire testing set and insert those predictions into the data set.

```
[ ]: X = test_df.drop(['pitcher_id', 'batter_id', 'stadium_id', 'umpire_id',  
→ 'catcher_id', 'is_strike', 'pitch_id'], axis=1)
```

```
X = pd.get_dummies(X, prefix=['pitcher', 'batter', 'is'],  
→ columns=['pitcher_side', 'batter_side', 'pitch_type'])
```

```
[ ]: predictions = tuned_xgb.predict(X)
```

```
[ ]: test_df['is_strike'] = predictions  
test_df.head(10)
```

```
[ ]: test_df.to_csv('../Data/predicted_test_set.csv')
```

4 Conclusion

Overall we created a good model to predict if a specific pitch was going to be a strike or not. We started off by cleaning the data set and making sure that no null values were in the table. Next we had to create the target variable from a column that already existed in the data set. Finally had to wrangle some data to get it all in the correct format to be suitable to run a machine learning model on. Using gradient boosting, tuned a number of hyper-parameters, and made predictions full training set. After making predictions, we needed to make sure the model was performing well, and took a look at a number of different metrics for model performance. Accuracy, the classification report, AUC score, and Matthew’s Correlation Coefficient all agreed that this model we built was well suited for predicting strikes. We also took a look at the most important features of the model, and came away with plate side and plate height being the most important features.

If I had more time and resources to dedicate to this project, I would have tuned some more of the XGBoost hyper-parameters to make this model even more accurate. Having only tuned five parameters, there could be some more room for improvement, but the time it could have taken to do so may have outweighed the gains produce by finding more optimal parameters.