

# Mariners Challenge

February 19, 2020

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline

import xgboost as xgb
from xgboost import XGBClassifier

from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report, precision_recall_curve, auc, \
    matthews_corrcoef, accuracy_score

from joblib import load
```

```
[2]: train_df = pd.read_csv('../Data/2020-train.csv')
```

```
[3]: train_df.head()
```

```
[3]:  pitcher_id pitcher_side batter_id batter_side stadium_id umpire_id \
0    d7e3acce          Right  32678d8d          Right   a4833794  f88d09f4
1    44ec1bf5          Right   81d51733          Left   f60d6ea5  b67d862c
2    44d87ee6          Left   8eefccb7          Right   a9b8b538  13993d26
3    ff6adae0          Right   8f8ab5af          Right   e569ec39  0d8ba4bb
4    c70c96e5          Right  10874746          Right   a5ce1bf6  94a4c552

   catcher_id  inning  top_bottom  outs  ...  zone_speed  vert_approach_angle \
0    83cdf9ff        3           1   0.0  ...   86.024200           -4.37258
1    a126f66f        6           2   0.0  ...   89.458199           -4.90467
2    9db4e46f        5           2   2.0  ...   75.593597           -6.00728
3    bbbfd290        5           1   2.0  ...   76.396400           -9.50640
4    75087ec8        8           1   2.0  ...   83.215302           -4.53233

   horz_approach_angle  zone_time      x55  y55      z55 pitch_type \
0           1.429580    0.404622 -0.059343   55   6.03322          FA
1           -2.148410    0.385719 -2.148680   55   6.23380          FA
2           -0.122044    0.463953  1.300450   55   6.14750          CH
```

|   |           |          |           |    |         |    |
|---|-----------|----------|-----------|----|---------|----|
| 3 | -2.581980 | 0.458471 | -1.659590 | 55 | 6.60043 | CU |
| 4 | -0.268188 | 0.415965 | -1.526170 | 55 | 4.77332 | FA |

|   | pitch_call   | pitch_id |
|---|--------------|----------|
| 0 | InPlay       | 42fce2f6 |
| 1 | InPlay       | 3e9cda86 |
| 2 | BallCalled   | f129a6cd |
| 3 | InPlay       | 03e9bc05 |
| 4 | StrikeCalled | 48feb675 |

[5 rows x 36 columns]

```
[4]: train_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 582205 entries, 0 to 582204
Data columns (total 36 columns):
pitcher_id          582205 non-null object
pitcher_side        582205 non-null object
batter_id           582205 non-null object
batter_side         582205 non-null object
stadium_id          582205 non-null object
umpire_id           582205 non-null object
catcher_id          582205 non-null object
inning              582205 non-null int64
top_bottom          582205 non-null int64
outs                582053 non-null float64
balls               582205 non-null int64
strikes             582205 non-null int64
release_speed       582093 non-null float64
vert_release_angle  582093 non-null float64
horz_release_angle  582093 non-null float64
spin_rate           573194 non-null float64
spin_axis           582093 non-null float64
tilt                580953 non-null object
rel_height          582093 non-null float64
rel_side            582093 non-null float64
extension           582093 non-null float64
vert_break          582093 non-null float64
induced_vert_break  582093 non-null float64
horz_break          582093 non-null float64
plate_height        582139 non-null float64
plate_side          582139 non-null float64
zone_speed          582093 non-null float64
vert_approach_angle 582093 non-null float64
horz_approach_angle 582093 non-null float64
zone_time           582093 non-null float64
```

```

x55          582093 non-null float64
y55          582205 non-null int64
z55          582093 non-null float64
pitch_type   581720 non-null object
pitch_call   582205 non-null object
pitch_id     582205 non-null object
dtypes: float64(20), int64(5), object(11)
memory usage: 159.9+ MB

```

## 1 Data Cleaning

First let's get rid of the null values floating around in the data set. There are a total of 582,205 entries in the initial data set, some columns are full but some columns have null values in them. Starting with release speed, it does not seem like there are many null values, let's take a look.

```
[5]: train_df[train_df['release_speed'].isnull()]
```

```
[5]:
```

|        | pitcher_id | pitcher_side | batter_id | batter_side | stadium_id | umpire_id | \ |
|--------|------------|--------------|-----------|-------------|------------|-----------|---|
| 3405   | 8759809c   | Right        | 00bde845  | Right       | 80756f45   | 9c6cbb5e  |   |
| 12334  | f6d227a5   | Right        | 69426d29  | Left        | a3f610ed   | 9c6cbb5e  |   |
| 13895  | b74a40d9   | Right        | 00bde845  | Right       | a3f610ed   | 9c6cbb5e  |   |
| 16260  | 47032f76   | Left         | 20bf9444  | Right       | c9712626   | 9c6cbb5e  |   |
| 17008  | 96a1cebe   | Right        | 76c0475e  | Right       | a5ce1bf6   | 9c6cbb5e  |   |
| ...    | ...        | ...          | ...       | ...         | ...        | ...       |   |
| 564267 | b3336756   | Left         | bbbfd290  | Right       | 9b5daeaf   | 9c6cbb5e  |   |
| 564950 | c3ededfb   | Right        | 8f8ab5af  | Right       | b20853fa   | 9c6cbb5e  |   |
| 567947 | ad6bf2a7   | Right        | 066e327d  | Right       | fe6b0f40   | 9c6cbb5e  |   |
| 571936 | 0a606b2d   | Right        | 336a9f05  | Left        | 1a39a252   | 9c6cbb5e  |   |
| 573697 | f556cf83   | Left         | 30300714  | Right       | b20853fa   | 9c6cbb5e  |   |

|        | catcher_id | inning | top_bottom | outs | ... | zone_speed | \ |
|--------|------------|--------|------------|------|-----|------------|---|
| 3405   | 9c6cbb5e   | 5      | 2          | 0.0  | ... | NaN        |   |
| 12334  | 9c6cbb5e   | 5      | 2          | 2.0  | ... | NaN        |   |
| 13895  | 9c6cbb5e   | 3      | 1          | 1.0  | ... | NaN        |   |
| 16260  | 9c6cbb5e   | 16     | 2          | 1.0  | ... | NaN        |   |
| 17008  | 9c6cbb5e   | 9      | 2          | 0.0  | ... | NaN        |   |
| ...    | ...        | ...    | ...        | ...  | ... | ...        |   |
| 564267 | 9c6cbb5e   | 9      | 1          | 2.0  | ... | NaN        |   |
| 564950 | 9c6cbb5e   | 1      | 1          | 2.0  | ... | NaN        |   |
| 567947 | 9c6cbb5e   | 5      | 2          | 0.0  | ... | NaN        |   |
| 571936 | 9c6cbb5e   | 9      | 1          | 1.0  | ... | NaN        |   |
| 573697 | 9c6cbb5e   | 7      | 1          | 0.0  | ... | NaN        |   |

|       | vert_approach_angle | horz_approach_angle | zone_time | x55 | y55 | z55 | \ |
|-------|---------------------|---------------------|-----------|-----|-----|-----|---|
| 3405  | NaN                 | NaN                 | NaN       | NaN | 55  | NaN |   |
| 12334 | NaN                 | NaN                 | NaN       | NaN | 55  | NaN |   |
| 13895 | NaN                 | NaN                 | NaN       | NaN | 55  | NaN |   |

|        |     |     |     |     |     |     |
|--------|-----|-----|-----|-----|-----|-----|
| 16260  | NaN | NaN | NaN | NaN | 55  | NaN |
| 17008  | NaN | NaN | NaN | NaN | 55  | NaN |
| ...    | ... | ... | ... | ... | ... | ... |
| 564267 | NaN | NaN | NaN | NaN | 55  | NaN |
| 564950 | NaN | NaN | NaN | NaN | 55  | NaN |
| 567947 | NaN | NaN | NaN | NaN | 55  | NaN |
| 571936 | NaN | NaN | NaN | NaN | 55  | NaN |
| 573697 | NaN | NaN | NaN | NaN | 55  | NaN |

|        | pitch_type | pitch_call   | pitch_id |
|--------|------------|--------------|----------|
| 3405   | NaN        | StrikeCalled | 1d66612a |
| 12334  | NaN        | InPlay       | beb842a8 |
| 13895  | NaN        | InPlay       | e271ef9d |
| 16260  | FA         | InPlay       | 49f26761 |
| 17008  | SL         | BallCalled   | ca8c6341 |
| ...    | ...        | ...          | ...      |
| 564267 | CH         | InPlay       | 8d21a585 |
| 564950 | NaN        | BallCalled   | 142a06b3 |
| 567947 | NaN        | StrikeCalled | 5326a5f6 |
| 571936 | NaN        | BallCalled   | 3c4e2fe4 |
| 573697 | NaN        | StrikeCalled | 763a95f4 |

[112 rows x 36 columns]

Only 112, I feel comfortable dropping these and it not affecting the integrity of the data.

```
[6]: train_df = train_df.drop(train_df[train_df['release_speed'].isnull()].index)
train_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 582093 entries, 0 to 582204
Data columns (total 36 columns):
pitcher_id          582093 non-null object
pitcher_side        582093 non-null object
batter_id           582093 non-null object
batter_side         582093 non-null object
stadium_id          582093 non-null object
umpire_id           582093 non-null object
catcher_id          582093 non-null object
inning              582093 non-null int64
top_bottom          582093 non-null int64
outs                581941 non-null float64
balls               582093 non-null int64
strikes             582093 non-null int64
release_speed       582093 non-null float64
vert_release_angle  582093 non-null float64
horz_release_angle  582093 non-null float64
spin_rate           573194 non-null float64
```

```

spin_axis          582093 non-null float64
tilt               580953 non-null object
rel_height         582093 non-null float64
rel_side           582093 non-null float64
extension          582093 non-null float64
vert_break         582093 non-null float64
induced_vert_break 582093 non-null float64
horz_break         582093 non-null float64
plate_height       582093 non-null float64
plate_side         582093 non-null float64
zone_speed         582093 non-null float64
vert_approach_angle 582093 non-null float64
horz_approach_angle 582093 non-null float64
zone_time          582093 non-null float64
x55                582093 non-null float64
y55                582093 non-null int64
z55                582093 non-null float64
pitch_type         581674 non-null object
pitch_call         582093 non-null object
pitch_id           582093 non-null object
dtypes: float64(20), int64(5), object(11)
memory usage: 164.3+ MB

```

Now let's take a look at outs. Again, not many null values in it, let's see how many there are.

```
[7]: train_df[train_df['outs'].isnull()]
```

```

[7]:   pitcher_id pitcher_side batter_id batter_side stadium_id umpire_id \
6689    a6118212      Right  5dce2d1c      Right   99faafae  c16da957
14410    5b740fab      Left  e2e2a336      Right   d0e0eb76  c9752165
19161    161160dd      Right f338e9d3      Left    1a39a252  c229ef9e
23843    f6d227a5      Right e9553a98      Right   d0e0eb76  26a1bb6b
27156    264562c6      Right 29d12af7      Left    a5ce1bf6  9806dfbc
...      ...      ...      ...      ...      ...
559476    b3d5c0a9      Right cf690f2f      Left    0c59f5af  8f1ef267
569804    91700130      Right de9d396f      Left    99faafae  667d5752
569836    91700130      Right de9d396f      Left    99faafae  667d5752
570351    d5ef78cb      Right 781ec6be      Right   aa998b21  a9ad7586
574039    b3d5c0a9      Right cf690f2f      Left    0c59f5af  8f1ef267

      catcher_id  inning  top_bottom  outs  ...  zone_speed  \
6689    fd37f21c      7           1   NaN  ...   75.911797
14410    5a42193e      5           2   NaN  ...   72.887398
19161    41ac8158      4           2   NaN  ...   84.766296
23843    0ffec018      6           1   NaN  ...   74.321999
27156    a3a2988b     11           2   NaN  ...   86.122299
...      ...      ...      ...      ...      ...
559476    fa18ff59      5           2   NaN  ...   80.375298

```

|        |          |   |   |     |     |           |
|--------|----------|---|---|-----|-----|-----------|
| 569804 | 41ac8158 | 7 | 1 | NaN | ... | 84.243103 |
| 569836 | 41ac8158 | 7 | 1 | NaN | ... | 82.846298 |
| 570351 | a3a2988b | 6 | 1 | NaN | ... | 85.266098 |
| 574039 | fa18ff59 | 5 | 2 | NaN | ... | 67.233902 |

|        | vert_approach_angle | horz_approach_angle | zone_time | x55      | y55 | \ |
|--------|---------------------|---------------------|-----------|----------|-----|---|
| 6689   | -10.27540           | -1.215380           | 0.457509  | -1.44079 | 55  |   |
| 14410  | -7.79636            | 0.764407            | 0.476133  | 3.30888  | 55  |   |
| 19161  | -5.21754            | 0.053372            | 0.417402  | -1.30256 | 55  |   |
| 23843  | -9.18817            | -4.715870           | 0.470084  | -2.66638 | 55  |   |
| 27156  | -7.29354            | -4.722650           | 0.411571  | -2.93787 | 55  |   |
| ...    | ...                 | ...                 | ...       | ...      | ... |   |
| 559476 | -7.07432            | -2.223690           | 0.430519  | -1.15777 | 55  |   |
| 569804 | -5.38618            | 0.594491            | 0.419565  | -1.64532 | 55  |   |
| 569836 | -6.36660            | -1.029200           | 0.425654  | -1.62191 | 55  |   |
| 570351 | -4.60246            | -1.696110           | 0.401914  | -1.92533 | 55  |   |
| 574039 | -11.05840           | -2.251130           | 0.515505  | -1.19462 | 55  |   |

|        | z55     | pitch_type | pitch_call     | pitch_id |
|--------|---------|------------|----------------|----------|
| 6689   | 6.19782 | SL         | StrikeSwinging | 8f7e287a |
| 14410  | 5.70064 | CH         | BallCalled     | fdaab946 |
| 19161  | 5.76697 | FA         | BallCalled     | 479ee407 |
| 23843  | 6.70092 | CH         | StrikeSwinging | b01234ef |
| 27156  | 6.37265 | SL         | StrikeSwinging | f7771432 |
| ...    | ...     | ...        | ...            | ...      |
| 559476 | 6.08150 | SL         | FoulBall       | d29aa13d |
| 569804 | 6.31180 | FA         | BallCalled     | d711100f |
| 569836 | 6.24978 | FA         | StrikeCalled   | 5f05e674 |
| 570351 | 5.80506 | FA         | StrikeSwinging | 24ac5f78 |
| 574039 | 5.88421 | CU         | FoulBall       | 3a28f2f5 |

[152 rows x 36 columns]

Only 152, that's good. I'll drop these as well since there's no real way to know how many outs there were on a given pitch.

```
[8]: train_df = train_df.drop(train_df[train_df['outs'].isnull()].index)
train_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 581941 entries, 0 to 582204
Data columns (total 36 columns):
pitcher_id      581941 non-null object
pitcher_side    581941 non-null object
batter_id       581941 non-null object
batter_side     581941 non-null object
stadium_id      581941 non-null object
umpire_id       581941 non-null object
```

```

catcher_id          581941 non-null object
inning              581941 non-null int64
top_bottom          581941 non-null int64
outs                581941 non-null float64
balls               581941 non-null int64
strikes             581941 non-null int64
release_speed       581941 non-null float64
vert_release_angle  581941 non-null float64
horz_release_angle  581941 non-null float64
spin_rate           573044 non-null float64
spin_axis           581941 non-null float64
tilt                580802 non-null object
rel_height          581941 non-null float64
rel_side            581941 non-null float64
extension           581941 non-null float64
vert_break          581941 non-null float64
induced_vert_break  581941 non-null float64
horz_break          581941 non-null float64
plate_height        581941 non-null float64
plate_side          581941 non-null float64
zone_speed          581941 non-null float64
vert_approach_angle 581941 non-null float64
horz_approach_angle 581941 non-null float64
zone_time           581941 non-null float64
x55                 581941 non-null float64
y55                 581941 non-null int64
z55                 581941 non-null float64
pitch_type          581522 non-null object
pitch_call          581941 non-null object
pitch_id            581941 non-null object
dtypes: float64(20), int64(5), object(11)
memory usage: 164.3+ MB

```

Now let's take a look at spin rate. It certainly looks like there are a lot of data points missing in this column, so dropping all the values may not be the best idea. Let's see how many there are.

```
[9]: train_df[train_df['spin_rate'].isnull()]
```

```

[9]:   pitcher_id pitcher_side batter_id batter_side stadium_id umpire_id \
43      4c807a49      Left  210e8d5b      Right  402559d3  4ff102e5
148     cb113772     Right  96339e13      Left  aa998b21  c683b9a6
161     af6d3149     Right  6c43d395     Right  03722f5d  d057fd71
237     eccb6087     Right  08b0b39d      Left  a5ce1bf6  9806dfbc
304     e332e67d     Right  34a8f234      Left  402559d3  fbbea103
...      ...      ...      ...      ...      ...
581856  00f5fb90     Right  b4efd4bf      Right  f682daed  cac8185e
581859  fa9b0925     Right  c1ec06e6     Right  5025d8df  7675ce83
581927  09da5d7a     Left  073c2b16     Right  b20853fa  4db7bcbc

```

|        |          |       |          |       |          |          |
|--------|----------|-------|----------|-------|----------|----------|
| 581992 | be5181f0 | Right | 566220c7 | Right | 03722f5d | a86853a2 |
| 582189 | a2f05755 | Right | e7a70ed1 | Left  | 854c6c72 | 16750c18 |

|        | catcher_id | inning | top_bottom | outs | ... | zone_speed | \ |
|--------|------------|--------|------------|------|-----|------------|---|
| 43     | fbc0970f   | 6      | 2          | 2.0  | ... | 77.869301  |   |
| 148    | a3a2988b   | 7      | 1          | 0.0  | ... | 73.582497  |   |
| 161    | e9aa50df   | 9      | 1          | 1.0  | ... | 76.742798  |   |
| 237    | b1499101   | 9      | 1          | 1.0  | ... | 79.280602  |   |
| 304    | e4fac104   | 5      | 2          | 2.0  | ... | 80.649803  |   |
| ...    | ...        | ...    | ...        | ...  | ... | ...        |   |
| 581856 | e4fac104   | 4      | 2          | 1.0  | ... | 77.819801  |   |
| 581859 | 054f7d9f   | 6      | 1          | 0.0  | ... | 77.702797  |   |
| 581927 | 9db4e46f   | 9      | 1          | 0.0  | ... | 76.172997  |   |
| 581992 | 5b8927f6   | 2      | 2          | 2.0  | ... | 77.229103  |   |
| 582189 | daa1322d   | 4      | 2          | 2.0  | ... | 77.375603  |   |

|        | vert_approach_angle | horz_approach_angle | zone_time | x55      | y55 | \ |
|--------|---------------------|---------------------|-----------|----------|-----|---|
| 43     | -6.96547            | 2.47110             | 0.444589  | 2.32704  | 55  |   |
| 148    | -9.61581            | -2.45899            | 0.474079  | -1.42864 | 55  |   |
| 161    | -6.86927            | -3.62681            | 0.461172  | -3.17556 | 55  |   |
| 237    | -6.75395            | -3.49224            | 0.438489  | -1.96139 | 55  |   |
| 304    | -8.90885            | -2.33627            | 0.433380  | -1.18638 | 55  |   |
| ...    | ...                 | ...                 | ...       | ...      | ... |   |
| 581856 | -8.53157            | -2.76430            | 0.446821  | -2.01818 | 55  |   |
| 581859 | -10.84580           | -3.94847            | 0.452332  | -1.45562 | 55  |   |
| 581927 | -8.79066            | 3.64529             | 0.459917  | 2.00739  | 55  |   |
| 581992 | -7.42066            | -3.82928            | 0.458138  | -1.49327 | 55  |   |
| 582189 | -8.75378            | -2.46304            | 0.450004  | -1.84355 | 55  |   |


|        | z55     | pitch_type | pitch_call     | pitch_id |
|--------|---------|------------|----------------|----------|
| 43     | 5.75515 | SL         | InPlay         | 1769b4d5 |
| 148    | 5.84820 | SL         | BallCalled     | ef2a6b0e |
| 161    | 4.88357 | SL         | StrikeCalled   | f31fc865 |
| 237    | 5.75645 | SL         | StrikeCalled   | 3588e119 |
| 304    | 6.42062 | SL         | InPlay         | 6fa4b697 |
| ...    | ...     | ...        | ...            | ...      |
| 581856 | 5.79492 | SL         | InPlay         | fc151903 |
| 581859 | 6.02164 | CU         | BallCalled     | 941e91ca |
| 581927 | 6.51264 | SL         | StrikeSwinging | 3f7dd184 |
| 581992 | 5.40118 | SL         | StrikeSwinging | 40cc471a |
| 582189 | 5.49470 | SL         | InPlay         | 059d09ff |

[8897 rows x 36 columns]

Almost 9000 rows. Dropping that many data points could have an adverse effect on the data set and any potential modeling we do. Let's take a look at a countplot and see the distribution of the spin rate column.



```
[10]: plt.figure(figsize=(12,9))
sns.distplot(train_df[train_df['spin_rate'].notnull()]['spin_rate'], hist=True,
             kde=False)
plt.xlabel('Spin rate')
plt.ylabel('Count')
plt.title('Count of spin rates')
plt.show()
```



It looks like a normal shaped curve, which is expected from this type of data. Filling the missing values with the average would be the best way to go. That keeps the normal distribution of this variable intact and shouldn't have any adverse effect on any future modeling.

```
[11]: train_df['spin_rate'].describe()
```

```
[11]: count    573044.000000
      mean      2220.693335
      std       311.989506
      min       362.382996
      25%      2072.879883
      50%      2238.449951
      75%      2400.790039
```

```
max          3752.239990
Name: spin_rate, dtype: float64
```

```
[12]: train_df['spin_rate'].fillna(train_df['spin_rate'].mean(), inplace=True)
```

```
[13]: train_df['spin_rate'].describe()
```

```
[13]: count      581941.000000
      mean        2220.693335
      std         309.595392
      min         362.382996
      25%        2076.080078
      50%        2233.979980
      75%        2397.649902
      max         3752.239990
      Name: spin_rate, dtype: float64
```

Tilt is the next column to have missing values, let's take a look at how many.

```
[14]: train_df[train_df['tilt'].isnull()]
```

```
[14]:
```

|        | pitcher_id | pitcher_side | batter_id | batter_side | stadium_id | umpire_id | \ |
|--------|------------|--------------|-----------|-------------|------------|-----------|---|
| 378    | cd483725   | Right        | 192899a6  | Right       | c9712626   | 1869cf54  |   |
| 770    | 22b76a09   | Left         | d11080ae  | Left        | 5025d8df   | 598ea1b2  |   |
| 1339   | cd483725   | Right        | deb2ab32  | Right       | 0b15e1ca   | eb059a22  |   |
| 2117   | 28e273c4   | Left         | 4ac005f3  | Right       | d0d69f32   | 46051258  |   |
| 3399   | 193d153f   | Left         | ed874f19  | Right       | c9712626   | eb059a22  |   |
| ...    | ...        | ...          | ...       | ...         | ...        | ...       |   |
| 580815 | 98eaf8b2   | Right        | f57085ec  | Right       | 99faafae   | 373947e5  |   |
| 581637 | 44ec1bf5   | R            | b3dac04c  | L           | cfe02944   | aea4dd5a  |   |
| 581716 | 60a6f8df   | Right        | cddcbd8f  | Left        | 0faa3b2d   | 4581c636  |   |
| 581733 | f45c0602   | Right        | ad84b429  | Left        | fe6b0f40   | bb04ea23  |   |
| 582161 | 1fb18290   | Left         | e14059d7  | Right       | 6a69d99b   | ff7406e8  |   |

|        | catcher_id | inning | top_bottom | outs | ... | zone_speed | \ |
|--------|------------|--------|------------|------|-----|------------|---|
| 378    | 00ae6fb5   | 3      | 1          | 1.0  | ... | 73.954399  |   |
| 770    | a421b54b   | 3      | 2          | 0.0  | ... | 67.967903  |   |
| 1339   | 00ae6fb5   | 6      | 2          | 1.0  | ... | 75.172897  |   |
| 2117   | 4fedda83   | 3      | 1          | 1.0  | ... | 75.375900  |   |
| 3399   | 4f9cd7f9   | 4      | 1          | 2.0  | ... | 75.718102  |   |
| ...    | ...        | ...    | ...        | ...  | ... | ...        |   |
| 580815 | ccd72da8   | 9      | 2          | 2.0  | ... | 79.889999  |   |
| 581637 | a126f66f   | 6      | 2          | 1.0  | ... | 88.644798  |   |
| 581716 | a421b54b   | 7      | 1          | 0.0  | ... | 79.617599  |   |
| 581733 | dc18f830   | 1      | 1          | 2.0  | ... | 82.575500  |   |
| 582161 | bbbfd290   | 8      | 2          | 2.0  | ... | 71.934700  |   |

|        | vert_approach_angle | horz_approach_angle | zone_time | x55       | y55 | \ |
|--------|---------------------|---------------------|-----------|-----------|-----|---|
| 378    | -9.60881            | -1.961030           | 0.477100  | -1.181070 | 55  |   |
| 770    | -11.73220           | 1.604870            | 0.512340  | 1.140470  | 55  |   |
| 1339   | -9.79158            | -1.867760           | 0.471512  | -0.749778 | 55  |   |
| 2117   | -8.49879            | 2.576980            | 0.459569  | 2.623400  | 55  |   |
| 3399   | -9.39808            | 2.573850            | 0.462214  | 2.150780  | 55  |   |
| ...    | ...                 | ...                 | ...       | ...       | ... |   |
| 580815 | -8.93767            | -1.194220           | 0.435881  | -1.759100 | 55  |   |
| 581637 | -5.31458            | -2.812660           | 0.390148  | -2.398420 | 55  |   |
| 581716 | -6.18068            | -2.975590           | 0.435664  | -2.043060 | 55  |   |
| 581733 | -5.96087            | -0.656144           | 0.416289  | -2.912520 | 55  |   |
| 582161 | -13.77550           | 1.224220            | 0.486536  | 1.447500  | 55  |   |

|        | z55     | pitch_type | pitch_call     | pitch_id |
|--------|---------|------------|----------------|----------|
| 378    | 6.13692 | CU         | InPlay         | d30f5214 |
| 770    | 6.32628 | CU         | BallCalled     | a7209bc8 |
| 1339   | 5.86303 | CU         | StrikeSwinging | 6b05ddc4 |
| 2117   | 6.40245 | CU         | InPlay         | 3face29a |
| 3399   | 6.07629 | CU         | StrikeSwinging | e1c3703d |
| ...    | ...     | ...        | ...            | ...      |
| 580815 | 6.48687 | CU         | StrikeSwinging | 672aa57b |
| 581637 | 6.25742 | FA         | BallCalled     | 9e8fb97a |
| 581716 | 5.01476 | SL         | StrikeCalled   | 1d36d6f5 |
| 581733 | 5.64645 | CH         | FoulBall       | 2e8f759b |
| 582161 | 6.16835 | CU         | StrikeSwinging | 7db0fd25 |

[1139 rows x 36 columns]

1139 rows with null values. That's not much in the grand scheme of the data set, only about 0.2% of the entire data set. Dropping these would not hurt in the long run.

```
[15]: train_df = train_df.drop(train_df[train_df['tilt'].isnull()].index)
```

```
[16]: train_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 580802 entries, 0 to 582204
Data columns (total 36 columns):
pitcher_id      580802 non-null object
pitcher_side    580802 non-null object
batter_id       580802 non-null object
batter_side     580802 non-null object
stadium_id      580802 non-null object
umpire_id       580802 non-null object
catcher_id      580802 non-null object
inning          580802 non-null int64
top_bottom      580802 non-null int64
outs            580802 non-null float64
```

```

balls                580802 non-null int64
strikes              580802 non-null int64
release_speed        580802 non-null float64
vert_release_angle   580802 non-null float64
horz_release_angle   580802 non-null float64
spin_rate            580802 non-null float64
spin_axis            580802 non-null float64
tilt                 580802 non-null object
rel_height           580802 non-null float64
rel_side             580802 non-null float64
extension            580802 non-null float64
vert_break           580802 non-null float64
induced_vert_break   580802 non-null float64
horz_break           580802 non-null float64
plate_height         580802 non-null float64
plate_side           580802 non-null float64
zone_speed           580802 non-null float64
vert_approach_angle  580802 non-null float64
horz_approach_angle  580802 non-null float64
zone_time            580802 non-null float64
x55                  580802 non-null float64
y55                  580802 non-null int64
z55                  580802 non-null float64
pitch_type           580383 non-null object
pitch_call           580802 non-null object
pitch_id             580802 non-null object
dtypes: float64(20), int64(5), object(11)
memory usage: 164.0+ MB

```

Finally, `pitch_type` is the last column to have null values in it. Let's take a look.

```
[17]: train_df[train_df['pitch_type'].isnull()]
```

```

[17]:   pitcher_id pitcher_side batter_id batter_side stadium_id umpire_id \
634      bff0f759      Left  192899a6      Right  d0d69f32  9c6cbb5e
3126     900e6090     Right  699983d6      Left  78aaa563  9c6cbb5e
3263     7bdd4794     Right  44924919     Right  83508f28  9c6cbb5e
3485     7bdd4794     Right  0b8c61b3      Left  83508f28  9c6cbb5e
4655     57613174     Right  00bde845     Right  78aaa563  9c6cbb5e
...      ...      ...      ...      ...      ...
576312    b48cf592     Right  a3b17b9b      Left  cfe02944  9c6cbb5e
578590    b4eadd6d     Right  0ae0de45     Right  0a0cfe0d  9c6cbb5e
579026    57613174     Right  6b115fe9      Left  fe6b0f40  9c6cbb5e
579999    d629b647     Right  699983d6      Left  78aaa563  9c6cbb5e
580965    57613174     Right  fd347bb1      Left  78aaa563  9c6cbb5e

      catcher_id  inning  top_bottom  outs  ...  zone_speed  \
634      9c6cbb5e       5           2    0.0  ...    67.892998

```

|        |          |     |     |     |     |           |
|--------|----------|-----|-----|-----|-----|-----------|
| 3126   | 9c6cbb5e | 8   | 2   | 1.0 | ... | 81.340797 |
| 3263   | 9c6cbb5e | 9   | 1   | 0.0 | ... | 86.714600 |
| 3485   | 9c6cbb5e | 8   | 1   | 1.0 | ... | 85.167503 |
| 4655   | 9c6cbb5e | 12  | 1   | 1.0 | ... | 85.837502 |
| ...    | ...      | ... | ... | ... | ... | ...       |
| 576312 | 9c6cbb5e | 8   | 2   | 1.0 | ... | 74.434799 |
| 578590 | 9c6cbb5e | 8   | 1   | 1.0 | ... | 87.362602 |
| 579026 | 9c6cbb5e | 8   | 2   | 0.0 | ... | 77.284798 |
| 579999 | 9c6cbb5e | 14  | 2   | 2.0 | ... | 83.999802 |
| 580965 | 9c6cbb5e | 14  | 1   | 0.0 | ... | 86.906998 |

|        | vert_approach_angle | horz_approach_angle | zone_time | x55      | y55 | \   |
|--------|---------------------|---------------------|-----------|----------|-----|-----|
| 634    | -7.08524            | 0.721203            | 0.516315  | 2.17234  | 55  |     |
| 3126   | -5.03602            | -2.062410           | 0.432663  | -2.91834 | 55  |     |
| 3263   | -3.29493            | -1.750810           | 0.402108  | -1.73126 | 55  |     |
| 3485   | -4.46310            | -1.934850           | 0.410087  | -1.83128 | 55  |     |
| 4655   | -4.32100            | -1.389630           | 0.402448  | -1.74871 | 55  |     |
| ...    | ...                 | ...                 | ...       | ...      | ... | ... |
| 576312 | -1.24151            | -1.315220           | 0.465865  | -1.59554 | 55  |     |
| 578590 | -4.12619            | -1.554950           | 0.396194  | -1.33279 | 55  |     |
| 579026 | -7.36536            | -1.514690           | 0.460128  | -1.91246 | 55  |     |
| 579999 | -5.26501            | 0.381528            | 0.411825  | -1.65040 | 55  |     |
| 580965 | -3.75225            | -0.884140           | 0.397061  | -1.94958 | 55  |     |

|        | z55     | pitch_type | pitch_call   | pitch_id |
|--------|---------|------------|--------------|----------|
| 634    | 6.50402 | NaN        | StrikeCalled | 34257ee5 |
| 3126   | 5.70209 | NaN        | InPlay       | 972b06d9 |
| 3263   | 5.63128 | NaN        | BallCalled   | 44187be2 |
| 3485   | 5.61829 | NaN        | StrikeCalled | bf482788 |
| 4655   | 5.99225 | NaN        | InPlay       | f41ccb06 |
| ...    | ...     | ...        | ...          | ...      |
| 576312 | 1.38558 | NaN        | FoulBall     | 17b5318a |
| 578590 | 6.22610 | NaN        | FoulBall     | d20f5e52 |
| 579026 | 5.90440 | NaN        | FoulBall     | 528ecba1 |
| 579999 | 6.96849 | NaN        | FoulBall     | e0d8dea9 |
| 580965 | 5.92467 | NaN        | InPlay       | fba61f04 |

[419 rows x 36 columns]

Only 419 rows, these can be dropped.

```
[18]: train_df = train_df.drop(train_df[train_df['pitch_type'].isnull()].index)
train_df = train_df.reset_index().drop('index', axis=1)
```

```
[19]: train_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 580383 entries, 0 to 580382
```

```

Data columns (total 36 columns):
pitcher_id          580383 non-null object
pitcher_side        580383 non-null object
batter_id           580383 non-null object
batter_side         580383 non-null object
stadium_id          580383 non-null object
umpire_id           580383 non-null object
catcher_id          580383 non-null object
inning              580383 non-null int64
top_bottom          580383 non-null int64
outs                580383 non-null float64
balls               580383 non-null int64
strikes             580383 non-null int64
release_speed       580383 non-null float64
vert_release_angle  580383 non-null float64
horz_release_angle  580383 non-null float64
spin_rate           580383 non-null float64
spin_axis           580383 non-null float64
tilt                580383 non-null object
rel_height          580383 non-null float64
rel_side            580383 non-null float64
extension           580383 non-null float64
vert_break          580383 non-null float64
induced_vert_break  580383 non-null float64
horz_break          580383 non-null float64
plate_height        580383 non-null float64
plate_side          580383 non-null float64
zone_speed          580383 non-null float64
vert_approach_angle 580383 non-null float64
horz_approach_angle 580383 non-null float64
zone_time           580383 non-null float64
x55                 580383 non-null float64
y55                 580383 non-null int64
z55                 580383 non-null float64
pitch_type          580383 non-null object
pitch_call          580383 non-null object
pitch_id            580383 non-null object
dtypes: float64(20), int64(5), object(11)
memory usage: 159.4+ MB

```

```
[20]: train_df.head(10)
```

```

[20]:  pitcher_id pitcher_side batter_id batter_side stadium_id umpire_id \
0    d7e3acce          Right  32678d8d          Right   a4833794   f88d09f4
1    44ec1bf5          Right  81d51733          Left    f60d6ea5   b67d862c
2    44d87ee6          Left   8eefccb7          Right   a9b8b538   13993d26
3    ff6adae0          Right  8f8ab5af          Right   e569ec39   0d8ba4bb

```

|   |          |       |          |       |          |          |
|---|----------|-------|----------|-------|----------|----------|
| 4 | c70c96e5 | Right | 10874746 | Right | a5ce1bf6 | 94a4c552 |
| 5 | 98f8936a | Right | a58e31f3 | Right | 9b5daeaf | 0dad94e8 |
| 6 | 28e273c4 | Left  | 9a2db1f2 | Right | d0d69f32 | caf1f50b |
| 7 | 4f3062b6 | Left  | 7e2bb9a9 | Right | c9712626 | 33bb973b |
| 8 | afae9816 | Left  | ffe7832e | Left  | d0d69f32 | f88d09f4 |
| 9 | 61ab8c67 | Right | daa1322d | Right | f682daed | c4c41d26 |

|   | catcher_id | inning | top_bottom | outs | ... | zone_speed | vert_approach_angle | \ |
|---|------------|--------|------------|------|-----|------------|---------------------|---|
| 0 | 83cdf9ff   | 3      | 1          | 0.0  | ... | 86.024200  | -4.37258            |   |
| 1 | a126f66f   | 6      | 2          | 0.0  | ... | 89.458199  | -4.90467            |   |
| 2 | 9db4e46f   | 5      | 2          | 2.0  | ... | 75.593597  | -6.00728            |   |
| 3 | bbbfd290   | 5      | 1          | 2.0  | ... | 76.396400  | -9.50640            |   |
| 4 | 75087ec8   | 8      | 1          | 2.0  | ... | 83.215302  | -4.53233            |   |
| 5 | 68d1111a   | 7      | 2          | 0.0  | ... | 80.265404  | -8.24794            |   |
| 6 | 4fedda83   | 3      | 1          | 1.0  | ... | 87.948799  | -4.76645            |   |
| 7 | 20bf9444   | 6      | 1          | 1.0  | ... | 76.352798  | -10.25710           |   |
| 8 | 4fedda83   | 2      | 1          | 1.0  | ... | 78.281097  | -4.85101            |   |
| 9 | 41ac8158   | 4      | 1          | 1.0  | ... | 86.078400  | -6.09955            |   |

|   | horz_approach_angle | zone_time | x55       | y55 | z55     | pitch_type | \ |
|---|---------------------|-----------|-----------|-----|---------|------------|---|
| 0 | 1.429580            | 0.404622  | -0.059343 | 55  | 6.03322 | FA         |   |
| 1 | -2.148410           | 0.385719  | -2.148680 | 55  | 6.23380 | FA         |   |
| 2 | -0.122044           | 0.463953  | 1.300450  | 55  | 6.14750 | CH         |   |
| 3 | -2.581980           | 0.458471  | -1.659590 | 55  | 6.60043 | CU         |   |
| 4 | -0.268188           | 0.415965  | -1.526170 | 55  | 4.77332 | FA         |   |
| 5 | 0.780148            | 0.438111  | -2.075230 | 55  | 5.79080 | CH         |   |
| 6 | 0.696210            | 0.390590  | 2.569990  | 55  | 6.09316 | FA         |   |
| 7 | 4.681720            | 0.464021  | 2.497290  | 55  | 6.22659 | SL         |   |
| 8 | 1.945590            | 0.440157  | 2.559950  | 55  | 5.91159 | FA         |   |
| 9 | 0.425454            | 0.411268  | -0.876224 | 55  | 6.53540 | FA         |   |

|   | pitch_call     | pitch_id |
|---|----------------|----------|
| 0 | InPlay         | 42fce2f6 |
| 1 | InPlay         | 3e9cda86 |
| 2 | BallCalled     | f129a6cd |
| 3 | InPlay         | 03e9bc05 |
| 4 | StrikeCalled   | 48feb675 |
| 5 | StrikeSwinging | 419540c7 |
| 6 | StrikeSwinging | cf85249d |
| 7 | BallCalled     | c9423da3 |
| 8 | FoulBall       | 51ad39b4 |
| 9 | FoulBall       | b89e4ec3 |

[10 rows x 36 columns]

No more null values in the data set now, now we can move on to creating the target variable for modeling.

## 2 Data Wrangling

In the test set, the target variable is called “is\_strike”, and we don’t have a column like that here in the training set. However, we do have a “pitch\_call” column, which we can use to create the “is\_strike” column. Along with called strikes and swinging strikes, any ball batted in to play or any foul balls are also counted as strikes. Using this, we can build the “is\_strike” column using a simple for loop.

```
[21]: train_df['pitch_call'].unique()
```

```
[21]: array(['InPlay', 'BallCalled', 'StrikeCalled', 'StrikeSwinging',  
          'FoulBall', 'HitByPitch', 'BallIntentional'], dtype=object)
```

```
[22]: is_strike = []  
is_strike_list = ['InPlay', 'StrikeCalled', 'StrikeSwinging', 'FoulBall']  
for i in train_df['pitch_call']:  
    if i in is_strike_list:  
        is_strike.append(1)  
    else:  
        is_strike.append(0)
```

We can easily assign that list to a new column in the data set.

```
[23]: train_df['is_strike'] = is_strike
```

```
[24]: train_df.head(10)
```

```
[24]:  pitcher_id pitcher_side batter_id batter_side stadium_id umpire_id  \  
0    d7e3acce      Right  32678d8d      Right    a4833794  f88d09f4  
1    44ec1bf5      Right  81d51733      Left     f60d6ea5  b67d862c  
2    44d87ee6      Left   8eefccb7      Right    a9b8b538  13993d26  
3    ff6adae0      Right  8f8ab5af      Right    e569ec39  0d8ba4bb  
4    c70c96e5      Right  10874746      Right    a5ce1bf6  94a4c552  
5    98f8936a      Right  a58e31f3      Right    9b5daeaf  0dad94e8  
6    28e273c4      Left   9a2db1f2      Right    d0d69f32  caf1f50b  
7    4f3062b6      Left   7e2bb9a9      Right    c9712626  33bb973b  
8    afae9816      Left   ffe7832e      Left     d0d69f32  f88d09f4  
9    61ab8c67      Right  daa1322d      Right    f682daed  c4c41d26  
  
   catcher_id  inning  top_bottom  outs  ...  vert_approach_angle  \  
0    83cdf9ff       3           1   0.0  ...           -4.37258  
1    a126f66f       6           2   0.0  ...           -4.90467  
2    9db4e46f       5           2   2.0  ...           -6.00728  
3    bbbfd290       5           1   2.0  ...           -9.50640  
4    75087ec8       8           1   2.0  ...           -4.53233  
5    68d1111a       7           2   0.0  ...           -8.24794  
6    4fedda83       3           1   1.0  ...           -4.76645  
7    20bf9444       6           1   1.0  ...           -10.25710
```



|   |          |   |   |     |     |          |
|---|----------|---|---|-----|-----|----------|
| 8 | 4fedda83 | 2 | 1 | 1.0 | ... | -4.85101 |
| 9 | 41ac8158 | 4 | 1 | 1.0 | ... | -6.09955 |

|   | horz_approach_angle | zone_time | x55       | y55 | z55     | pitch_type \ |
|---|---------------------|-----------|-----------|-----|---------|--------------|
| 0 | 1.429580            | 0.404622  | -0.059343 | 55  | 6.03322 | FA           |
| 1 | -2.148410           | 0.385719  | -2.148680 | 55  | 6.23380 | FA           |
| 2 | -0.122044           | 0.463953  | 1.300450  | 55  | 6.14750 | CH           |
| 3 | -2.581980           | 0.458471  | -1.659590 | 55  | 6.60043 | CU           |
| 4 | -0.268188           | 0.415965  | -1.526170 | 55  | 4.77332 | FA           |
| 5 | 0.780148            | 0.438111  | -2.075230 | 55  | 5.79080 | CH           |
| 6 | 0.696210            | 0.390590  | 2.569990  | 55  | 6.09316 | FA           |
| 7 | 4.681720            | 0.464021  | 2.497290  | 55  | 6.22659 | SL           |
| 8 | 1.945590            | 0.440157  | 2.559950  | 55  | 5.91159 | FA           |
| 9 | 0.425454            | 0.411268  | -0.876224 | 55  | 6.53540 | FA           |

|   | pitch_call     | pitch_id | is_strike |
|---|----------------|----------|-----------|
| 0 | InPlay         | 42fce2f6 | 1         |
| 1 | InPlay         | 3e9cda86 | 1         |
| 2 | BallCalled     | f129a6cd | 0         |
| 3 | InPlay         | 03e9bc05 | 1         |
| 4 | StrikeCalled   | 48feb675 | 1         |
| 5 | StrikeSwinging | 419540c7 | 1         |
| 6 | StrikeSwinging | cf85249d | 1         |
| 7 | BallCalled     | c9423da3 | 0         |
| 8 | FoulBall       | 51ad39b4 | 1         |
| 9 | FoulBall       | b89e4ec3 | 1         |

[10 rows x 37 columns]

Now we have our data set with the target variable, let's take a look at the "is\_strike" column and its value counts.

```
[25]: train_df['is_strike'].value_counts()
```

```
[25]: 1    369807
      0    210576
      Name: is_strike, dtype: int64
```

Interesting. We have an imbalanced classification problem here, with the majority class being almost twice as large as the minority class. That has implications for modeling in the future, namely being careful about what classification model is used for this problem. We also may need to use some resampling methods if the model is choosing the majority class by an overwhelming margin.

Before we get into modeling however, there was something I noticed with the "tilt" column. It has two different types of string data packed into the column. We'll need to fix that column to get it all into one data format.

```
[26]: train_df['tilt'].unique()
```

```
[26]: array(['1:00', '12:15', '11:15', '7:45', '2:15', '2:45', '10:30', '4:45',  
          '11:00', '1:30', '1:15', '12:45', '5:15', '10:45', '6:30', '8:00',  
          '4:00', '2:00', '7:30', '3:30', '12:00', '1:45', '9:00', '10:00',  
          '11:30', '32400 secs', '12:30', '9:15', '11:45', '9:45', '10:15',  
          '3:00', '42300 secs', '5:00', '7:15', '7:00', '6:45', '9:30',  
          '8:45', '3:45', '43200 secs', '6:15', '2:30', '4:15', '35100 secs',  
          '5:45', '5:30', '8:15', '8:30', '29700 secs', '44100 secs', '3:15',  
          '4:30', '14400 secs', '6300 secs', '45000 secs', '38700 secs',  
          '34200 secs', '36000 secs', '45900 secs', '36900 secs',  
          '3600 secs', '18000 secs', '7200 secs', '40500 secs', '5400 secs',  
          '8100 secs', '15300 secs', '27900 secs', '4500 secs', '23400 secs',  
          '25200 secs', '41400 secs', '30600 secs', '6:00', '9900 secs',  
          '33300 secs', '37800 secs', '13500 secs', '27000 secs',  
          '39600 secs', '12600 secs', '17100 secs', '16200 secs',  
          '11700 secs', '9000 secs', '18900 secs', '26100 secs',  
          '22500 secs', '20700 secs', '24300 secs', '31500 secs',  
          '21600 secs', '10800 secs', '19800 secs', '28800 secs'],  
          dtype=object)
```

I'm choosing to turn all of the "1:00", "12:15" format into a seconds-from-midnight integer, that will be the easiest way to get all of the column into one data format and data type.

```
[27]: train_df['tilt'] = train_df['tilt'].map(lambda x: sum(a*int(t) for a, t in  
          ↪ zip([3600, 60], x.split(':')))) \n  
          if ':' in x else int(x[:-5]))
```

```
[28]: train_df.to_csv('../Data/model_data.csv')
```

### 3 Predictive Modeling

Now that our data set clean and how we want it, we can get into some predictive modeling. Seeing as this is a binary classification problem, we'll need to use a classification algorithm. I'm choosing to use gradient boosting here because I've used it in the past and have gotten good results with it in a timely manner. Logistic regression would be faster, but would give us a less accurate model than a gradient boosting model.

I'm going to do the hyperparameter tuning in its own dedicated notebook, then load the trained model into this notebook after it's been fitted with all the correct hyperparameters.

```
[29]: X = train_df.drop(['pitcher_id', 'batter_id', 'stadium_id', 'umpire_id',  
          ↪ 'catcher_id', 'pitch_call', 'is_strike', 'pitch_id'], axis=1)  
y = train_df['is_strike']  
  
X_train, X_test, y_train, y_test = train_test_split(X, y, stratify=y,  
          ↪ random_state=34)
```

Getting dummy variables for some of the categorical variables would be good for modeling. It will allow us to see how much importance was placed on these features by the model.

```
[30]: X_train = pd.get_dummies(X_train, prefix=['pitcher', 'batter', 'is'],
    ↳ columns=['pitcher_side', 'batter_side', 'pitch_type'])
X_test = pd.get_dummies(X_test, prefix=['pitcher', 'batter', 'is'],
    ↳ columns=['pitcher_side', 'batter_side', 'pitch_type'])
```

```
[31]: X_train.head()
```

```
[31]:
```

|        | inning | top_bottom | outs | balls | strikes | release_speed | \ |
|--------|--------|------------|------|-------|---------|---------------|---|
| 362341 | 8      | 1          | 1.0  | 3     | 2       | 86.056702     |   |
| 392255 | 3      | 2          | 0.0  | 0     | 1       | 93.787697     |   |
| 520345 | 3      | 1          | 0.0  | 0     | 0       | 90.838699     |   |
| 120374 | 3      | 1          | 1.0  | 1     | 1       | 89.893600     |   |
| 194486 | 7      | 1          | 1.0  | 0     | 0       | 94.105202     |   |

|        | vert_release_angle | horz_release_angle | spin_rate   | spin_axis   | ... | \ |
|--------|--------------------|--------------------|-------------|-------------|-----|---|
| 362341 | -0.754319          | -2.89459           | 2425.870117 | 205.197006  | ... |   |
| 392255 | -1.202140          | -3.24475           | 2496.909912 | 210.024994  | ... |   |
| 520345 | -1.730600          | 1.75061            | 2230.830078 | 158.481995  | ... |   |
| 120374 | -2.494800          | -3.26674           | 2175.179932 | -150.666000 | ... |   |
| 194486 | -2.047340          | -2.19327           | 2071.290039 | 227.212997  | ... |   |

|        | pitcher_Left | pitcher_Right | batter_Left | batter_Right | is_CH | is_CU | \ |
|--------|--------------|---------------|-------------|--------------|-------|-------|---|
| 362341 | 0            | 1             | 0           | 1            | 0     | 0     |   |
| 392255 | 0            | 1             | 1           | 0            | 0     | 0     |   |
| 520345 | 1            | 0             | 0           | 1            | 0     | 0     |   |
| 120374 | 0            | 1             | 0           | 1            | 0     | 0     |   |
| 194486 | 0            | 1             | 1           | 0            | 0     | 0     |   |

|        | is_FA | is_KN | is_SL | is_XX |
|--------|-------|-------|-------|-------|
| 362341 | 0     | 0     | 1     | 0     |
| 392255 | 1     | 0     | 0     | 0     |
| 520345 | 1     | 0     | 0     | 0     |
| 120374 | 0     | 0     | 1     | 0     |
| 194486 | 1     | 0     | 0     | 0     |

[5 rows x 36 columns]

Now that we have our modeling data set, let's get into modeling the data, first we will look at a non-tuned XGBoost model and see how it performs on the data, then compare that to the tuned model from the "Mariners Machine Learning Model" notebook.

```
[32]: base_xgb = XGBClassifier()
base_xgb.fit(X_train, y_train)
xgb_base_pred = base_xgb.predict(X_test)
```

```
[33]: print(f'Base XGB Classifier Test Accuracy: {round(accuracy_score(y_test,
    ↳xgb_base_pred) * 100, 2)}')
print('Base XGB Classifier Classification Report')
print(classification_report(y_test, xgb_base_pred))
print('\n')
print(f'Base XGB Classifier MCC Score: {matthews_corrcoef(y_test,
    ↳xgb_base_pred)}')

xgb_probs = base_xgb.predict_proba(X_test)
xgb_probs = xgb_probs[:, 1]

xgb_precision, xgb_recall, _ = precision_recall_curve(y_test, xgb_probs)

no_skill = len(y_test[y_test == 1]) / len(y_test)
plt.figure(figsize=(12,9))
plt.plot([0, 1], [no_skill, no_skill], linestyle='--', label='No Skill')
plt.plot(xgb_recall, xgb_precision, marker='.', label='Gradient Boosting')
plt.title('Precision Recall Curve for a Base XGB Classifier')
plt.xlabel('Recall')
plt.ylabel('Precision')
plt.legend()
plt.show()


print(f'Base XGB Classifier AUC Score: {auc(xgb_recall, xgb_precision)}')
```

Base XGB Classifier Test Accuracy: 86.46

Base XGB Classifier Classification Report

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.82      | 0.81   | 0.81     | 52644   |
| 1            | 0.89      | 0.90   | 0.89     | 92452   |
| accuracy     |           |        | 0.86     | 145096  |
| macro avg    | 0.85      | 0.85   | 0.85     | 145096  |
| weighted avg | 0.86      | 0.86   | 0.86     | 145096  |

Base XGB Classifier MCC Score: 0.706412788570753



Base XGB Classifier AUC Score: 0.9675354035506347

By itself, it's a good model. 86% accuracy and a solid precision/recall on both classes is a good start. As well as a .706 MCC score and very high AUC score. Looks like the base XGBoost is a good starting out point for the hyperparameter tuning we did in the other notebook. Let's load that in and take a look at the same metrics as above.

```
[34]: tuned_xgb = load('xgboost_model.pkl')
      xgb_tuned_pred = tuned_xgb.predict(X_test)

[35]: print(f'Tuned XGB Classifier Test Accuracy: {round(accuracy_score(y_test,
      ↪xgb_tuned_pred) * 100, 2)}%')
      print('Tuned XGB Classifier Classification Report')
      print(classification_report(y_test, xgb_tuned_pred))
      print('\n')
      print(f'Tuned XGB Classifier MCC Score: {matthews_corrcoef(y_test,
      ↪xgb_tuned_pred)}%')

      xgb_probs = tuned_xgb.predict_proba(X_test)
      xgb_probs = xgb_probs[:, 1]
```

```

xgb_precision, xgb_recall, _ = precision_recall_curve(y_test, xgb_probs)

no_skill = len(y_test[y_test == 1]) / len(y_test)
plt.figure(figsize=(12,9))
plt.plot([0, 1], [no_skill, no_skill], linestyle='--', label='No Skill')
plt.plot(xgb_recall, xgb_precision, marker='.', label='Gradient Boosting')
plt.title('Precision Recall Curve for a Tuned XGB Classifier')
plt.xlabel('Recall')
plt.ylabel('Precision')
plt.legend()
plt.show()

print(f'Tuned XGB Classifier AUC Score: {auc(xgb_recall, xgb_precision)}')

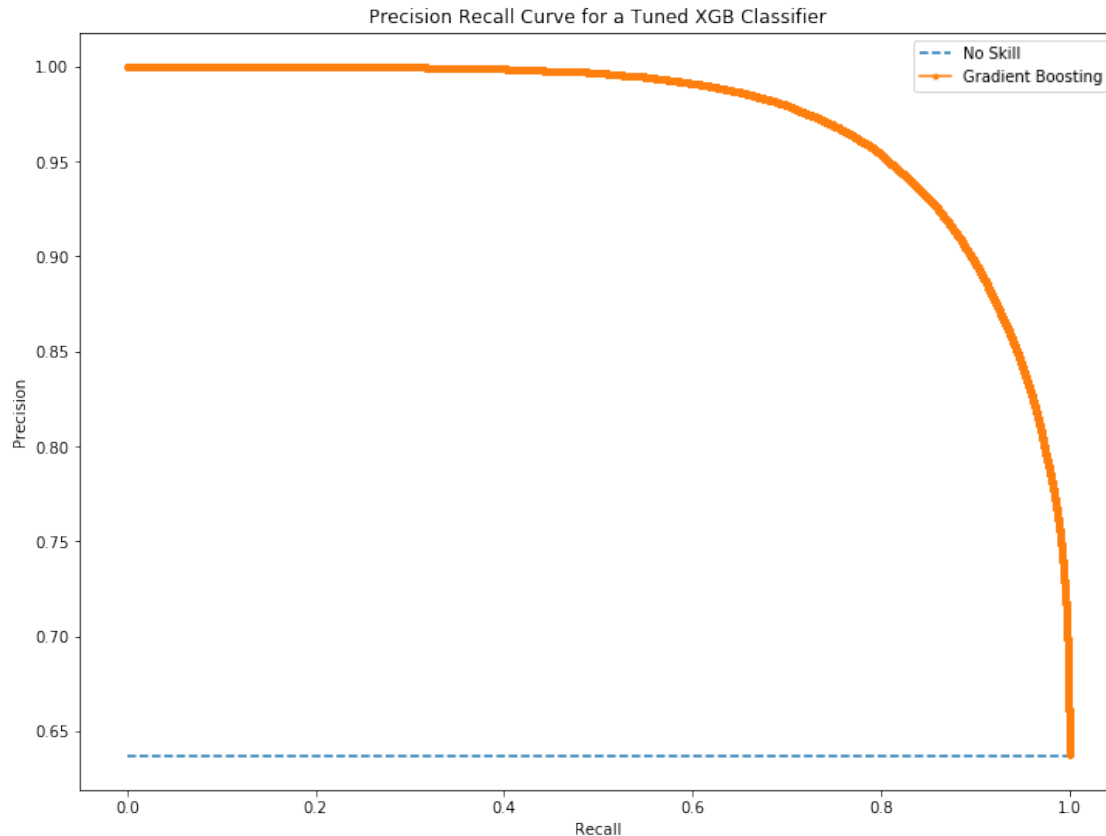
```

Tuned XGB Classifier Test Accuracy: 87.03

Tuned XGB Classifier Classification Report

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.82      | 0.82   | 0.82     | 52644   |
| 1            | 0.90      | 0.90   | 0.90     | 92452   |
| accuracy     |           |        | 0.87     | 145096  |
| macro avg    | 0.86      | 0.86   | 0.86     | 145096  |
| weighted avg | 0.87      | 0.87   | 0.87     | 145096  |

Tuned XGB Classifier MCC Score: 0.7190715044081059



Tuned XGB Classifier AUC Score: 0.9703777678129323

It did better, even if it was only slightly. Precision and recall scores on both classes improved, accuracy went up, and MCC and AUC scores went up as well. The hyper-parameter tuning we did in the other notebook worked well.

### 3.1 Making predictions on test set

We have a good model trained, now we can make predictions on the testing set. First we need to load it in and clean it the way we cleaned the training set.

```
[36]: test_df = pd.read_csv('../Data/2020-test.csv')
```

```
[37]: def clean_and_wrangle(df):  
  
    df = df.drop(df[df['release_speed'].isnull()].index)  
    df = df.drop(df[df['outs'].isnull()].index)  
    df['spin_rate'].fillna(df['spin_rate'].mean(), inplace=True)  
    df = df.drop(df[df['tilt'].isnull()].index)  
    df = df.drop(df[df['pitch_type'].isnull()].index)  
    df = df.reset_index().drop('index', axis=1)
```

```

df['tilt'] = df['tilt'].map(lambda x: sum(a*int(t) for a, t in zip([3600, 60], x.split(':')))) \
                                if ':' in x else int(x[:-5]))

return df

```

```
[38]: test_df = clean_and_wrangle(test_df)
```

```
[39]: test_df.head(10)
```

```
[39]:  pitcher_id pitcher_side batter_id batter_side stadium_id umpire_id \
0      d3396348          Left  d9b3bce2          Right  501b6728  a63083b5
1      4c807a49          Left  4aafd18a          Right  8d1f4cfc  93c9014b
2      18182a03          Right c790fbcb          Left   075be90a  9c02aab4
3      94a20652          Right bf921933          Right  934c75c6  043de890
4      4f3062b6          Left  65df5b42          Right  c9712626  d057fd71
5      3903adfd          Right 13448018          Left   45b7bf7c  1ce4b3e6
6      d9b3bce2          Right 1817dec7          Left   075be90a  0c8846f2
7      06e0842e          Left  730d2dbf          Left   20418ce9  852c6a22
8      fe5717f2          Right 44c206bb          Right  fe6b0f40  16750c18
9      7fa6b7cb          Right e5fe8773          Left   d0d69f32  c4c41d26

    catcher_id  inning  top_bottom  outs  ...  zone_speed  vert_approach_angle \
0      c338c856        8           1  2.0  ...   85.462196          -5.52951
1      97c420bc        1           2  1.0  ...   76.937698          -7.24994
2      568a8108        6           2  1.0  ...   83.710899          -7.12427
3      5e710b9e        5           1  1.0  ...   85.949799          -5.92277
4      00ae6fb5        3           1  1.0  ...   85.592598          -7.10051
5      fbc0970f        5           1  1.0  ...   83.406601          -5.14624
6      370c45c8        7           1  0.0  ...   84.818802          -4.38852
7      65b01821        6           2  1.0  ...   82.629601          -5.51211
8      62542678        6           1  2.0  ...   86.618500          -4.40207
9      9d29b427        5           2  0.0  ...   85.491203          -6.62900

    horz_approach_angle  zone_time    x55  y55    z55  pitch_type \
0           0.682682    0.411072  2.65519  55  6.30581          FA
1           0.617254    0.446388  2.05217  55  5.89617          CH
2          -4.845640    0.421318 -1.78613  55  5.93421          SL
3          -3.132810    0.400539 -1.56069  55  5.44192          FA
4           1.461540    0.406034  2.15070  55  6.43411          FA
5           0.550966    0.412397 -2.45849  55  6.32671          FA
6          -1.782670    0.405944 -2.91866  55  5.86711          FA
7           1.447130    0.423868  2.04815  55  5.84246          FA
8          -0.125200    0.397941 -1.50450  55  6.53886          FA
9          -0.512073    0.403582 -2.67956  55  5.61784          FA

    is_strike  pitch_id
0         NaN  f2204560

```



|   |     |          |
|---|-----|----------|
| 1 | NaN | 4a16102e |
| 2 | NaN | 73ffabd3 |
| 3 | NaN | 60ed54c3 |
| 4 | NaN | 5d720732 |
| 5 | NaN | 4aa772e1 |
| 6 | NaN | debae10c |
| 7 | NaN | c71b9d22 |
| 8 | NaN | fd77d5fb |
| 9 | NaN | c43cd8b2 |

[10 rows x 36 columns]

```
[40]: test_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 145097 entries, 0 to 145096
Data columns (total 36 columns):
pitcher_id          145097 non-null object
pitcher_side        145097 non-null object
batter_id           145097 non-null object
batter_side         145097 non-null object
stadium_id          145097 non-null object
umpire_id           145097 non-null object
catcher_id          145097 non-null object
inning              145097 non-null int64
top_bottom          145097 non-null int64
outs                145097 non-null float64
balls               145097 non-null int64
strikes             145097 non-null int64
release_speed       145097 non-null float64
vert_release_angle  145097 non-null float64
horz_release_angle  145097 non-null float64
spin_rate           145097 non-null float64
spin_axis           145097 non-null float64
tilt                145097 non-null int64
rel_height          145097 non-null float64
rel_side            145097 non-null float64
extension           145097 non-null float64
vert_break          145097 non-null float64
induced_vert_break  145097 non-null float64
horz_break          145097 non-null float64
plate_height        145097 non-null float64
plate_side          145097 non-null float64
zone_speed          145097 non-null float64
vert_approach_angle 145097 non-null float64
horz_approach_angle 145097 non-null float64
zone_time           145097 non-null float64
```

```

x55          145097 non-null float64
y55          145097 non-null int64
z55          145097 non-null float64
pitch_type   145097 non-null object
is_strike    0 non-null float64
pitch_id     145097 non-null object
dtypes: float64(21), int64(6), object(9)
memory usage: 39.9+ MB

```

No null values (except our target variable “is\_strike”). We’ll use our trained model from up above and make predictions on the entire testing set and insert those predictions into the data set.

```

[41]: X = test_df.drop(['pitcher_id', 'batter_id', 'stadium_id', 'umpire_id',
    ↪ 'catcher_id', 'is_strike', 'pitch_id'], axis=1)

X = pd.get_dummies(X, prefix=['pitcher', 'batter', 'is'],
    ↪ columns=['pitcher_side', 'batter_side', 'pitch_type'])

```

```

[42]: predictions = tuned_xgb.predict(X)

```

```

[43]: test_df['is_strike'] = predictions
test_df.head(10)

```

```

[43]:  pitcher_id pitcher_side batter_id batter_side stadium_id umpire_id  \
0    d3396348      Left  d9b3bce2      Right  501b6728  a63083b5
1    4c807a49      Left  4aafd18a      Right  8d1f4cfc  93c9014b
2    18182a03      Right c790fbef      Left  075be90a  9c02aab4
3    94a20652      Right bf921933      Right  934c75c6  043de890
4    4f3062b6      Left  65df5b42      Right  c9712626  d057fd71
5    3903adfd      Right 13448018      Left  45b7bf7c  1ce4b3e6
6    d9b3bce2      Right 1817dec7      Left  075be90a  0c8846f2
7    06e0842e      Left  730d2dbf      Left  20418ce9  852c6a22
8    fe5717f2      Right 44c206bb      Right  fe6b0f40  16750c18
9    7fa6b7cb      Right e5fe8773      Left  d0d69f32  c4c41d26

    catcher_id  inning  top_bottom  outs  ...  zone_speed  vert_approach_angle  \
0    c338c856      8          1  2.0  ...  85.462196      -5.52951
1    97c420bc      1          2  1.0  ...  76.937698      -7.24994
2    568a8108      6          2  1.0  ...  83.710899      -7.12427
3    5e710b9e      5          1  1.0  ...  85.949799      -5.92277
4    00ae6fb5      3          1  1.0  ...  85.592598      -7.10051
5    fbc0970f      5          1  1.0  ...  83.406601      -5.14624
6    370c45c8      7          1  0.0  ...  84.818802      -4.38852
7    65b01821      6          2  1.0  ...  82.629601      -5.51211
8    62542678      6          1  2.0  ...  86.618500      -4.40207
9    9d29b427      5          2  0.0  ...  85.491203      -6.62900

    horz_approach_angle  zone_time  x55  y55  z55  pitch_type  \

```

|   |           |          |          |    |         |    |
|---|-----------|----------|----------|----|---------|----|
| 0 | 0.682682  | 0.411072 | 2.65519  | 55 | 6.30581 | FA |
| 1 | 0.617254  | 0.446388 | 2.05217  | 55 | 5.89617 | CH |
| 2 | -4.845640 | 0.421318 | -1.78613 | 55 | 5.93421 | SL |
| 3 | -3.132810 | 0.400539 | -1.56069 | 55 | 5.44192 | FA |
| 4 | 1.461540  | 0.406034 | 2.15070  | 55 | 6.43411 | FA |
| 5 | 0.550966  | 0.412397 | -2.45849 | 55 | 6.32671 | FA |
| 6 | -1.782670 | 0.405944 | -2.91866 | 55 | 5.86711 | FA |
| 7 | 1.447130  | 0.423868 | 2.04815  | 55 | 5.84246 | FA |
| 8 | -0.125200 | 0.397941 | -1.50450 | 55 | 6.53886 | FA |
| 9 | -0.512073 | 0.403582 | -2.67956 | 55 | 5.61784 | FA |

|   | is_strike | pitch_id |
|---|-----------|----------|
| 0 | 0         | f2204560 |
| 1 | 1         | 4a16102e |
| 2 | 0         | 73ffabd3 |
| 3 | 0         | 60ed54c3 |
| 4 | 1         | 5d720732 |
| 5 | 0         | 4aa772e1 |
| 6 | 1         | debae10c |
| 7 | 1         | c71b9d22 |
| 8 | 1         | fd77d5fb |
| 9 | 1         | c43cd8b2 |

[10 rows x 36 columns]

## 4 Conclusion

Overall we created a good model to predict if a specific pitch was going to be a strike or not. We started off by cleaning the data set and making sure that no null values were in the table. Next we had to create the target variable from a column that already existed in the data set. Finally had to wrangle some data to get it all in the correct format to be suitable to run a machine learning model on. Using gradient boosting, tuned a number of hyper-parameters, and made predictions full training set. After making predictions, we needed to make sure the model was performing well, and took a look at a number of different metrics for model performance. Accuracy, the classification report, AUC score, and Matthew's Correlation Coefficient all agreed that this model we built was well suited for predicting strikes. We also took a look at the most important features of the model, and came away with plate side and plate height being the most important features.

If I had more time and resources to dedicate to this project, I would have tuned some more of the XGBoost hyper-parameters to make this model even more accurate. Having only tuned five parameters, there could be some more room for improvement, but the time it could have taken to do so may have outweighed the gains produce by finding more optimal parameters.