

Player Report

February 19, 2020

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import matplotlib.patches as patches
import seaborn as sns
%matplotlib inline

from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.metrics import confusion_matrix, classification_report, \
    precision_recall_curve, auc, matthews_corrcoef, accuracy_score
from sklearn.pipeline import Pipeline

from xgboost import XGBClassifier
```

```
[2]: train_df = pd.read_csv('../Data/2020-train.csv')
```

```
[3]: train_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 582205 entries, 0 to 582204
Data columns (total 36 columns):
pitcher_id          582205 non-null object
pitcher_side        582205 non-null object
batter_id           582205 non-null object
batter_side         582205 non-null object
stadium_id          582205 non-null object
umpire_id           582205 non-null object
catcher_id          582205 non-null object
inning              582205 non-null int64
top_bottom          582205 non-null int64
outs                582053 non-null float64
balls               582205 non-null int64
strikes             582205 non-null int64
release_speed       582093 non-null float64
vert_release_angle  582093 non-null float64
horz_release_angle  582093 non-null float64
spin_rate           573194 non-null float64
spin_axis           582093 non-null float64
```

```

tilt                580953 non-null object
rel_height          582093 non-null float64
rel_side            582093 non-null float64
extension           582093 non-null float64
vert_break          582093 non-null float64
induced_vert_break  582093 non-null float64
horz_break          582093 non-null float64
plate_height        582139 non-null float64
plate_side          582139 non-null float64
zone_speed          582093 non-null float64
vert_approach_angle 582093 non-null float64
horz_approach_angle 582093 non-null float64
zone_time           582093 non-null float64
x55                 582093 non-null float64
y55                 582205 non-null int64
z55                 582093 non-null float64
pitch_type          581720 non-null object
pitch_call          582205 non-null object
pitch_id            582205 non-null object
dtypes: float64(20), int64(5), object(11)
memory usage: 159.9+ MB

```

So we're looking at a specific catcher here and his pitch receiving skills, let's pick him out and see how big our data set is.

```
[4]: player_df = train_df[train_df['catcher_id'] == 'f06c9fdf']
player_df
```

```
[4]:
```

	pitcher_id	pitcher_side	batter_id	batter_side	stadium_id	umpire_id	\
35	f4414aa0	Left	f57085ec	Right	43dd6efb	a86853a2	
256	ed8aa16f	Right	fb0970f	Right	45b7bf7c	c229ef9e	
353	0d8afd1d	Left	1813cf1c	Left	d0e0eb76	9c02aab4	
397	870b67d9	Right	2d75ea3e	Right	43dd6efb	d057fd71	
418	f58fcfef	Right	277ee019	Right	d0e0eb76	5d49d16c	
...	
581483	309f3c2d	Right	675b0ce5	Left	d0e0eb76	5d49d16c	
581503	fe642525	Right	76c0475e	Right	a5ce1bf6	b0106fa5	
581550	12b0433b	Right	d59d0d36	Left	d0e0eb76	4e713444	
581864	d8352da3	Right	daa1322d	Right	934c75c6	51a1c7ee	
582073	309f3c2d	Right	0e9cc95b	Left	0c59f5af	3007964d	

	catcher_id	inning	top_bottom	outs	...	zone_speed	\
35	f06c9fdf	9	1	1.0	...	83.303299	
256	f06c9fdf	5	2	0.0	...	78.985100	
353	f06c9fdf	8	1	0.0	...	73.143700	
397	f06c9fdf	6	1	0.0	...	83.313301	
418	f06c9fdf	7	1	1.0	...	83.698898	
...	

581483	f06c9fdf	1	1	1.0	...	90.457703
581503	f06c9fdf	7	2	1.0	...	82.521301
581550	f06c9fdf	2	1	0.0	...	78.001503
581864	f06c9fdf	5	2	0.0	...	86.916801
582073	f06c9fdf	1	2	0.0	...	86.889099

	vert_approach_angle	horz_approach_angle	zone_time	x55	y55	\
35	-4.02502	0.834138	0.421826	0.759003	55	
256	-8.73613	-0.713829	0.451039	-1.810170	55	
353	-9.74545	4.512090	0.483157	2.432580	55	
397	-4.06607	-1.724440	0.413309	-2.126510	55	
418	-5.63199	0.283421	0.406953	-1.736550	55	
...	
581483	-4.08045	-1.345470	0.382682	-2.063990	55	
581503	-6.17842	-0.372579	0.417648	-2.032460	55	
581550	-6.18036	-0.723831	0.440129	-2.356510	55	
581864	-4.86654	1.099070	0.393080	-1.019420	55	
582073	-5.65405	-3.638330	0.397248	-2.655230	55	

	z55	pitch_type	pitch_call	pitch_id
35	6.17413	FA	BallCalled	c6607ace
256	6.29071	CH	FoulBall	8ef9aca4
353	6.46889	CU	FoulBall	b93024dc
397	6.00446	FA	BallCalled	2e4cf21e
418	5.69398	FA	FoulBall	b5157dae
...
581483	5.83140	FA	FoulBall	622c34e4
581503	5.92519	CH	StrikeCalled	ac39f15d
581550	6.00084	CH	StrikeSwinging	8040e585
581864	6.47430	FA	BallCalled	199f7047
582073	5.98179	FA	BallCalled	7f6b1f3b

[8822 rows x 36 columns]

8800 data points, not too bad. We can get some good data out of this. Let's see what happened on each of those pitches.


```
[5]: player_df['pitch_call'].value_counts()
```

```
[5]: BallCalled      3451
     InPlay          1516
     FoulBall        1460
     StrikeCalled    1381
     StrikeSwinging   980
     HitByPitch       34
     Name: pitch_call, dtype: int64
```

Lots of balls called, not good. Though that could be more on the pitcher that's throwing to him

than the catcher himself. Let's put these on a scatter plot to see where all of these pitches crossed the plate.


```
[6]: plt.figure(figsize=(12,9))
sns.scatterplot('plate_side', 'plate_height', data=player_df, hue='pitch_call')
plt.show()
```



That doesn't give a whole lot of information. Since we're looking at pitch receiving skills, let's take a look and called balls and strikes, perhaps we can see how well he's framing pitches.

```
[7]: framing_df = player_df[(player_df['pitch_call'] == 'StrikeCalled') |
    ↳ (player_df['pitch_call'] == 'BallCalled')]
```

```
[8]: plt.figure(figsize=(25,21))
sns.scatterplot('plate_side', 'plate_height', data=framing_df, hue='pitch_call')
plt.title('Balls or strikes called by location of pitch')
plt.show()
```



Interesting. It seems like this catcher is good at framing pitches, with the orange dots being strikes called extending into areas of majority blue dots. But, it also goes the other way around. Let's clean and wrangle our data set like we did in the "Mariners Challenge" notebook to run some machine learning on it and draw some conclusions about how well this catcher frames pitches.

```
[9]: def clean_and_wrangle(df):

    df = df.drop(df[df['release_speed'].isnull()].index)
    df = df.drop(df[df['outs'].isnull()].index)
    df['spin_rate'].fillna(df['spin_rate'].mean(), inplace=True)
    df = df.drop(df[df['tilt'].isnull()].index)
    df = df.drop(df[df['pitch_type'].isnull()].index)
    df = df.reset_index().drop('index', axis=1)

    df['tilt'] = df['tilt'].map(lambda x: sum(a*int(t) for a, t in zip([3600, 60], x.split(':')))) \
                                if ':' in x else int(x[:-5]))

    return df
```

```
[10]: framing_df = clean_and_wrangle(framing_df)
```

```
[11]: is_strike = []
      for i in framing_df['pitch_call']:
          if i == 'StrikeCalled':
              is_strike.append(1)
          else:
              is_strike.append(0)

      framing_df['is_strike'] = is_strike
```

```
[12]: framing_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4830 entries, 0 to 4829
Data columns (total 37 columns):
pitcher_id      4830 non-null object
pitcher_side    4830 non-null object
batter_id       4830 non-null object
batter_side     4830 non-null object
stadium_id      4830 non-null object
umpire_id       4830 non-null object
catcher_id      4830 non-null object
inning          4830 non-null int64
top_bottom      4830 non-null int64
outs            4830 non-null float64
balls           4830 non-null int64
strikes         4830 non-null int64
release_speed   4830 non-null float64
vert_release_angle 4830 non-null float64
horz_release_angle 4830 non-null float64
spin_rate       4830 non-null float64
spin_axis       4830 non-null float64
tilt            4830 non-null int64
rel_height      4830 non-null float64
rel_side        4830 non-null float64
extension       4830 non-null float64
vert_break      4830 non-null float64
induced_vert_break 4830 non-null float64
horz_break      4830 non-null float64
plate_height    4830 non-null float64
plate_side      4830 non-null float64
zone_speed      4830 non-null float64
vert_approach_angle 4830 non-null float64
horz_approach_angle 4830 non-null float64
zone_time       4830 non-null float64
x55             4830 non-null float64
```

```

y55          4830 non-null int64
z55          4830 non-null float64
pitch_type   4830 non-null object
pitch_call   4830 non-null object
pitch_id     4830 non-null object
is_strike    4830 non-null int64
dtypes: float64(20), int64(7), object(10)
memory usage: 1.4+ MB

```

After cleaning and getting our target variable, let's take a look at it and see if we're dealing with another imbalanced classification problem.

```
[13]: framing_df['is_strike'].value_counts()
```

```

[13]: 0    3450
      1    1380
      Name: is_strike, dtype: int64

```

Yes, a similar imbalanced classification problem to the “Mariners Challenge” notebook. We'll tune an XGB Classifier model to model our data and draw some conclusions from its predictions, let's get that model ready to go.

```
[14]: framing_df['pitch_type'].value_counts()
```

```

[14]: FA    2576
      SL    1029
      CH     694
      CU     528
      XX       3
      Name: pitch_type, dtype: int64

```

```

[15]: framing_df = framing_df.drop(framing_df[framing_df['pitch_type'] == 'XX'].index)
      framing_df = pd.get_dummies(framing_df, prefix=['pitcher', 'batter', 'is'],
      ↪columns=['pitcher_side', 'batter_side', 'pitch_type'])

```

```

[16]: X = framing_df.drop(['pitcher_id', 'batter_id', 'stadium_id', 'umpire_id',
      ↪'catcher_id', 'pitch_call', 'is_strike', 'pitch_id'], axis=1)
      y = framing_df['is_strike']

      X_train, X_test, y_train, y_test = train_test_split(X, y, stratify=y,
      ↪random_state=34, test_size=0.15)

```

```

[17]: steps = [('xgb', XGBClassifier(seed=34))]
      param_grid = {'xgb__n_estimators': np.arange(10, 210, 10),
      ↪'xgb__max_depth': np.arange(1,6,2)}
      pipeline = Pipeline(steps)
      cv_1 = GridSearchCV(pipeline, param_grid, cv=3)
      cv_1.fit(X_train, y_train)

```

```
print(cv_1.best_params_, cv_1.best_score_)
n_estimators = cv_1.best_params_['xgb__n_estimators']
max_depth = cv_1.best_params_['xgb__max_depth']
```

```
{'xgb__max_depth': 3, 'xgb__n_estimators': 150} 0.9261326933524986
```

```
[18]: xgb = XGBClassifier(n_estimators=n_estimators, max_depth=max_depth, seed=34)
xgb.fit(X_train, y_train)
xgb_pred = xgb.predict(X_test)
```

```
[19]: print(f'XGB Classifier Train Accuracy: {round(xgb.score(X_train, y_train) * 100, 2)}%')
print(f'XGB Classifier Test Accuracy: {round(xgb.score(X_test, y_test) * 100, 2)}%')
print('\n')
print('XGB Classifier Classification Report')
print(classification_report(y_test, xgb_pred))
print('\n')
print(f'XGB Classifier MCC Score: {matthews_corrcoef(y_test, xgb_pred)}')
```

```
XGB Classifier Train Accuracy: 95.95%
```

```
XGB Classifier Test Accuracy: 92.14%
```

```
XGB Classifier Classification Report
```

	precision	recall	f1-score	support
0	0.96	0.92	0.94	518
1	0.83	0.91	0.87	207
accuracy			0.92	725
macro avg	0.90	0.92	0.91	725
weighted avg	0.93	0.92	0.92	725

```
XGB Classifier MCC Score: 0.8149311048301918
```

Looks like a good model here. No signs of overfitting, it has a good recall and precision on both classes, and has an excellent MCC score. To evaluate this catcher's performance on framing pitches, I'm going to use a method for determining the number of "extra strikes" this catcher either gives up or gets called for him. I sourced it from this recent blog post from DataRobot: <https://blog.datarobot.com/catcher-pitch-framing-using-machine-learning-part-1>

The number of extra strikes is the number of false negatives - false positives. So the number of strikes the model thought were balls, minus the number of balls the model thought were strikes. This will tell us whether or not the catcher is getting more strikes called for his pitcher by framing the ball well. If the number is positive, that means the catcher is framing up the ball well and

getting more strikes called for his pitcher, and vice versa if the number is negative. Let's take a look.

```
[20]: compare_df = y_test.to_frame()
      compare_df['xgb_predict'] = xgb_pred
      compare_df
```

```
[20]:      is_strike  xgb_predict
      1149         0           0
      1537         0           0
      2543         0           0
      2933         0           0
      4762         0           0
      ...
      2731         1           1
      3951         0           0
      76         0           0
      4807         1           0
      2766         1           1
```

```
[725 rows x 2 columns]
```


```
[21]: len(compare_df[(compare_df['is_strike'] == 1) & (compare_df['xgb_predict'] == 0)]) - \
      len(compare_df[(compare_df['is_strike'] == 0) & (compare_df['xgb_predict'] == 1)])
```

```
[21]: -21
```

His extra strikes number is negative, which means he's not framing the ball very well and getting balls called that would normally be called for strikes. -21 isn't a huge number, so he's definitely not the worst at this, but he certainly has room for improvement. That was for the entire zone, so let's take a look at each side of the zone, left and right, and see what side he's better on. Starting with the right side.

```
[22]: right_framing = framing_df[framing_df['plate_side'] > 0]
```

```
[23]: plt.figure(figsize=(18,15))
      sns.scatterplot('plate_side', 'plate_height', data=right_framing, hue='pitch_call')
      plt.title('Balls or strikes called by location of pitch. Right side of zone.')
      plt.show()
```



```
[24]: X = right_framing.drop(['pitcher_id', 'batter_id', 'stadium_id', 'umpire_id',
    ↪ 'catcher_id', 'pitch_call', 'is_strike', 'pitch_id'], axis=1)
y = right_framing['is_strike']

X_train, X_test, y_train, y_test = train_test_split(X, y, stratify=y,
    ↪ random_state=34, test_size=0.15)
```

```
[25]: right_pred = xgb.predict(X_test)
```

```
[26]: right_compare_df = y_test.to_frame()
right_compare_df['right_predict'] = right_pred
right_compare_df
```

```
[26]:
```

	is_strike	right_predict
1859	1	1
504	0	0
2136	0	0
3876	0	0

1988	1	1
...
924	0	0
1113	0	0
2585	0	0
2448	0	0
940	1	1

[378 rows x 2 columns]


```
[27]: len(right_compare_df[(right_compare_df['is_strike'] == 1) &
    ↳(right_compare_df['right_predict'] == 0)]) - \
len(right_compare_df[(right_compare_df['is_strike'] == 0) &
    ↳(right_compare_df['right_predict'] == 1)])
```

[27]: -10

Still negative, which isn't good. Let's take a look at the left side and see if it's the same as the right side.

```
[28]: left_framing = framing_df[framing_df['plate_side'] < 0]
```

```
[29]: plt.figure(figsize=(18,15))
sns.scatterplot('plate_side', 'plate_height', data=left_framing,
    ↳hue='pitch_call')
plt.title('Balls or strikes called by location of pitch. Left side of zone.')
plt.show()
```



```
[30]: X = left_framing.drop(['pitcher_id', 'batter_id', 'stadium_id', 'umpire_id',
    ↪ 'catcher_id', 'pitch_call', 'is_strike', 'pitch_id'], axis=1)
y = left_framing['is_strike']

X_train, X_test, y_train, y_test = train_test_split(X, y, stratify=y,
    ↪ random_state=34, test_size=0.15)
```

```
[31]: left_pred = xgb.predict(X_test)
```

```
[32]: left_compare_df = y_test.to_frame()
left_compare_df['left_predict'] = left_pred
left_compare_df
```

```
[32]:
```

	is_strike	left_predict
616	1	0
746	1	1
3747	1	1
1019	0	0

513	0	0
...
981	1	1
1136	0	0
737	0	0
2587	0	0
886	0	0

[347 rows x 2 columns]


```
[33]: len(left_compare_df[(left_compare_df['is_strike'] == 1) &
    ↳ (left_compare_df['left_predict'] == 0)]) - \
len(left_compare_df[(left_compare_df['is_strike'] == 0) &
    ↳ (left_compare_df['left_predict'] == 1)])
```

[33]: -5

Less negative, which means this catcher does a slightly better job framing up pitches on the left side of the zone than the right side of the zone. What about from righty or lefty pitchers? Let's take a look at those data points and see who he's better for framing the ball for.

```
[34]: rp_framing = framing_df[framing_df['pitcher_Right'] == 1]
```

```
[35]: plt.figure(figsize=(18,15))
sns.scatterplot('plate_side', 'plate_height', data=rp_framing, hue='pitch_call')
plt.title('Balls or strikes called by location of pitch. Left side of zone.')
plt.show()
```



```
[36]: X = rp_framing.drop(['pitcher_id', 'batter_id', 'stadium_id', 'umpire_id',
    ↳ 'catcher_id', 'pitch_call', 'is_strike', 'pitch_id'], axis=1)
y = rp_framing['is_strike']

X_train, X_test, y_train, y_test = train_test_split(X, y, stratify=y,
    ↳ random_state=34, test_size=0.15)
```

```
[37]: rp_pred = xgb.predict(X_test)
```

```
[38]: rp_compare_df = y_test.to_frame()
rp_compare_df['rp_predict'] = rp_pred
rp_compare_df
```

```
[38]:
```

	is_strike	rp_predict
312	1	1
4551	0	0
2985	1	1
1986	0	0

1426	1	1
...
1031	0	0
2287	0	0
1250	1	1
2041	1	1
4802	0	0

[477 rows x 2 columns]


```
[39]: len(rp_compare_df[(rp_compare_df['is_strike'] == 1) &
    ↳ (rp_compare_df['rp_predict'] == 0)]) - \
len(rp_compare_df[(rp_compare_df['is_strike'] == 0) &
    ↳ (rp_compare_df['rp_predict'] == 1)])
```

[39]: -9

Similar number to the right side of the zone, let's take a look at the lefties.

```
[40]: lp_framing = framing_df[framing_df['pitcher_Left'] == 1]
```

```
[41]: plt.figure(figsize=(18,15))
sns.scatterplot('plate_side', 'plate_height', data=lp_framing, hue='pitch_call')
plt.title('Balls or strikes called by location of pitch. Left handed pitchers.')
plt.show()
```



```
[42]: X = lp_framing.drop(['pitcher_id', 'batter_id', 'stadium_id', 'umpire_id',
    ↳ 'catcher_id', 'pitch_call', 'is_strike', 'pitch_id'], axis=1)
y = lp_framing['is_strike']

X_train, X_test, y_train, y_test = train_test_split(X, y, stratify=y,
    ↳ random_state=34, test_size=0.15)
```

```
[43]: rp_pred = xgb.predict(X_test)
```

```
[44]: rp_compare_df = y_test.to_frame()
rp_compare_df['rp_predict'] = rp_pred
rp_compare_df
```

```
[44]:
```

	is_strike	rp_predict
98	0	0
4522	0	0
1661	0	0
913	0	0

2131	0	0
...
4703	0	0
1038	0	0
1788	0	0
2208	0	0
34	1	1

[248 rows x 2 columns]

```
[45]: len(rp_compare_df[(rp_compare_df['is_strike'] == 1) &
    ↳ (rp_compare_df['rp_predict'] == 0)]) - \
len(rp_compare_df[(rp_compare_df['is_strike'] == 0) &
    ↳ (rp_compare_df['rp_predict'] == 1)])
```

[45]: -4

Smaller number, similar to the left side of the zone. This catcher is better at framing pitches on the left side of the zone from left handed pitchers. He could be good to bring in for lefty specialist type pitchers, for throwing inside to left handed batters or low and away from right handed batters. Overall not a bad catcher, but definitely has room for improvement.